

Recap

- ▶ We have been considering linear classifiers:

$$\begin{aligned} h(X) &= 1 \text{ if } \sum_{i=1}^{d'} w_i \phi_i(X) + w_0 > 0 \\ &= 0 \text{ Otherwise} \end{aligned}$$

where ϕ_i are fixed functions.

- ▶ We take $h(X) = \text{sign}(W^T \Phi(X))$ for simplicity of notation. (Augmented feature vector).
- ▶ Perceptron is a classical algorithm to learn such a classifier.

Recap

- ▶ We also discussed linear regression. The objective is to learn a model:

$$\hat{y}(X) = W^T X$$

(We could use $\phi_i(X)$ in place of x_i).

- ▶ The least squares criterion is to minimize

$$J(W) = \frac{1}{2} \sum_{i=1}^n (W^T X_i - y_i)^2$$

- ▶ The minimizer is the linear least squares solution

$$W^* = (A^T A)^{-1} A^T Y$$

Where A is a matrix whose rows are X_i^T and Y is a vector whose components are y_i .

Recap

- ▶ We can also minimize J by iterative gradient descent.
- ▶ A stochastic gradient descent on J is the LMS algorithm.
- ▶ The LMS algorithm is:

$$W(k+1) = W(k) - \eta (X(k)^T W(k) - y(k)) X(k)$$

where $(X(k), y(k))$ is the (random) sample picked and $W(k)$ is the weight vector at iteration k .

- ▶ This is used in many adaptive signal processing problems.

Recap

- ▶ We can model the conditional distribution of y given X as Gaussian with mean $W^T X$.
- ▶ Then the ML estimate for W is same as the least squares solution.
- ▶ Thus, we can think of least squares criterion as appropriate when the measurements are corrupted by additive gaussian noise.

- ▶ The Gaussian noise assumption is alright for a regression problem.
- ▶ In a 2-class classification problem, where $y \in \{0, 1\}$, Gaussian noise does not make sense.
- ▶ So, we will investigate a slightly different model for the classification problem.

- ▶ We showed that if we want to predict y as a function of X to minimize $E[(f(X) - y)^2]$, then the optimal function is

$$f^*(X) = E[y | X]$$

- ▶ Suppose $y \in \{0, 1\}$. Then

$$f^*(X) = E[y | X] = P[y = 1 | X] = q_1(X)$$

- ▶ Hence in the least squares method we are learning an approximation to posterior probability.
- ▶ We can ask what would be a reasonable model for posterior probability.

By Bayes rule

$$\begin{aligned} q_0(X) &= \frac{f_0(X) p_0}{f_0(X) p_0 + f_1(X) p_1} \\ &= \frac{1}{1 + \frac{f_1(X) p_1}{f_0(X) p_0}} \\ &= \frac{1}{1 + \exp(-\xi)} \quad \text{where} \end{aligned}$$

$$\xi = -\ln \left(\frac{f_1(X) p_1}{f_0(X) p_0} \right) = \ln \left(\frac{f_0(X) p_0}{f_1(X) p_1} \right)$$

- ▶ Suppose f_0, f_1 are Gaussian with equal covariance matrices.
- ▶ Then

$$\ln \left(\frac{f_0(X) p_0}{f_1(X) p_1} \right) = W^T X + w_0$$

where

$$W = \Sigma^{-1}(\mu_0 - \mu_1); \quad w_0 = \frac{1}{2}(\mu_1^T \Sigma^{-1} \mu_1 - \mu_0^T \Sigma^{-1} \mu_0) + \ln \left(\frac{p_0}{p_1} \right)$$


- ▶ Thus, in this case we have

$$q_0(X) = f_{y|X}(0|X) = \frac{1}{1 + \exp(-W^T X - w_0)} = \sigma(W^T X + w_0)$$

This is called the logistic function or the sigmoid function.

Logistic Function

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



Logistic Regression

- ▶ The logistic function is a good model for posterior probability.
- ▶ We want to learn a probability model:

$$f_{y|X}(1|X) = P[y = 1|X] = \frac{1}{1 + \exp(-W^T X)}$$

(where we are assuming augmented feature vector)

- ▶ This is a discriminative model.
- ▶ We can learn the model using maximum likelihood approach
- ▶ Such a method is called logistic regression.

Logistic Regression

- ▶ Let the data be $\mathcal{D} = \{(X_1, y_1), \dots, (X_n, y_n)\}$.
- ▶ We are modelling y conditioned on X as Bernoulli with parameter $\sigma(W^T X)$.
- ▶ Hence the likelihood function is given by

$$L(W | \mathcal{D}) = \prod_{i=1}^n \theta_i^{y_i} (1 - \theta_i)^{1-y_i}, \quad \theta_i = \sigma(W^T X_i)$$

- ▶ Hence the log likelihood is given by

$$l(W | \mathcal{D}) = \sum_{i=1}^n [y_i \ln(\sigma(W^T X_i)) + (1 - y_i) \ln(1 - \sigma(W^T X_i))]$$

- ▶ To maximize this, we need its gradient.
- ▶ Note that the derivative of logistic function is

$$\frac{d\sigma(x)}{dx} = \sigma(x)(1 - \sigma(x))$$

- ▶ The log likelihood is given by

$$l(W|\mathcal{D}) = \sum_{i=1}^n [y_i \ln(\sigma(W^T X_i)) + (1 - y_i) \ln(1 - \sigma(W^T X_i))]$$

- ▶ Its gradient is

$$\begin{aligned} \frac{\partial l}{\partial W} &= \sum_{i=1}^n \left[y_i \frac{\sigma(W^T X_i)(1 - \sigma(W^T X_i))X_i}{\sigma(W^T X_i)} + \right. \\ &\quad \left. (1 - y_i) \frac{(-1)\sigma(W^T X_i)(1 - \sigma(W^T X_i))X_i}{1 - \sigma(W^T X_i)} \right] \\ &= \sum_{i=1}^n [y_i(1 - \sigma(W^T X_i))X_i + (1 - y_i)(-1)\sigma(W^T X_i)X_i] \\ &= \sum_{i=1}^n [(y_i - y_i\sigma(W^T X_i) - \sigma(W^T X_i) + y_i\sigma(W^T X_i)) X_i] \\ &= \sum_{i=1}^n [(y_i - \sigma(W^T X_i)) X_i] \end{aligned}$$

- ▶ In logistic regression we need to maximize this log likelihood.
- ▶ Hence the algorithm is

$$W(k+1) = W(k) + \eta \sum_{i=1}^n [(y_i - \sigma(W^T X_i)) X_i]$$

- ▶ The least squares algorithm to minimize $\sum_{i=1}^n (y_i - W^T X_i)^2$ through gradient descent is

$$W(k+1) = W(k) + \eta \sum_{i=1}^n [(y_i - W^T X_i) X_i]$$

- ▶ So, in logistic regression we are simply taking the model output as $\sigma(W^T X_i)$.
- ▶ In Least squares we would threshold $W^T X$ at 0.5 to classify a new X .
- ▶ In logistic regression, we threshold $\sigma(W^T X)$ at 0.5 to classify a new X .

- ▶ Since we know that logistic function is a good model for posterior probability we could (alternately) try and minimize

$$J(W) = \sum_{i=1}^n (y_i - \sigma(W^T X_i))^2$$

- ▶ Gradient descent on this would give us

$$W(k+1) = W(k) + \eta \sum_{i=1}^n [(y_i - \sigma(W^T X_i)) g(W, X_i) X_i]$$

where $g(W, X) = \sigma(W^T X_i)(1 - \sigma(W^T X_i))$.

- ▶ Very similar to the logistic regression algorithm.

- ▶ We are fitting a model $f(X) = W^T \Phi(X)$ to the data.
- ▶ We want to rate different W for their 'goodness of fit'.
- ▶ $\sum_i (W^T \Phi(X_i) - y_i)^2$ is the 'data error'.
- ▶ But it does not tell whole story of how good is W .

- ▶ We do not want the algorithm to 'overfit'.
- ▶ That is, low training error but high test error or generalization error
- ▶ Overfitting often happens if the model we learnt is 'too complex'.
- ▶ Hence we should have an algorithm that prefers 'simple' models.
- ▶ One way of doing this is through the so called regularization.

Regularization

- ▶ We can change our criterion to

$$\begin{aligned} J(W) &= \text{Data error} + \lambda \text{ model complexity} \\ &= \frac{1}{2} \sum_{i=1}^n (W^T \Phi(X_i) - y_i)^2 + \lambda \Omega(W) \end{aligned}$$

- ▶ Here $\Omega(W)$ is some measure of how 'complex' the model is.
- ▶ This is called regularized least squares and λ is called the regularization constant.

Ridge Regression

- ▶ In linear least squares regression, we often choose $\Omega(W) = \frac{1}{2}||W||^2$. Called Tikhonov regularization.
- ▶ Now the criterion is

$$\begin{aligned} J(W) &= \frac{1}{2} \sum_{i=1}^n (W^T \Phi(X_i) - y_i)^2 + \frac{\lambda}{2} W^T W \\ &= \frac{1}{2} (AW - Y)^T (AW - Y) + \frac{\lambda}{2} W^T W \end{aligned}$$

where, as earlier, A is the matrix whose rows are $\Phi(X_i)$.

- ▶ Equating the gradient of J to zero, we get

$$A^T (AW - Y) + \lambda W = 0$$

- ▶ This gives us

$$(A^T A + \lambda I)W = A^T Y \Rightarrow W^* = (A^T A + \lambda I)^{-1} A^T Y$$

- ▶ $A^T A$ is positive semidefinite
($z^T A^T A z = (Az)^T (Az) \geq 0$).
- ▶ Hence, $(A^T A + \lambda I)$ would have all eigen values above λ and hence is always invertible.
- ▶ So, regularization helps the condition number of the linear equations and thus the learning is more 'stable'.

- ▶ In ridge regression we are minimizing

$$J(W) = \frac{1}{2} \sum_{i=1}^n (W^T X_i - y_i)^2 + \frac{\lambda}{2} W^T W$$

- ▶ This is essentially same as solving the problem

$$\min_W \sum_{i=1}^n (W^T X_i - y_i)^2, \quad s.t. \quad \|W\| \leq B$$

(for an appropriate choice of λ).

- ▶ Here the norm used is the L_2 norm.
- ▶ So, ridge regression is said to use L_2 penalty.

- ▶ We could use other norms, e.g., L_1 norm.
- ▶ Thus we could reformulate this as minimizing

$$J(W) = \frac{1}{2} \sum_{i=1}^n (W^T X_i - y_i)^2 + \frac{\lambda}{2} \|W\|_1$$

where $\|W\|_1 = \sum_{i=1}^d |w_i|$.

- ▶ This formulation is called Lasso:

$$\min_W \sum_{i=1}^n (W^T X_i - y_i)^2, \quad s.t. \quad \|W\|_1 \leq B$$

- ▶ This is also a convex optimization problem.
- ▶ It promotes sparsity in the solution and hence is also suitable for 'feature selection'.

- ▶ Another way to look at regularized least squares is from a Bayesian framework.
- ▶ We saw that the least squares solution can also be derived as a ML estimate of parameters of a (reasonable) probability model for conditional distribution of y given X .
- ▶ The regularized least squares can be derived as a Bayesian (MAP) estimate of the parameters of the same model.

- ▶ As earlier, take the probability model as

$$f(y \mid X, W, \sigma) = \frac{1}{\sigma \sqrt{2\pi}} \exp \left(-\frac{1}{2} \frac{(y - W^T X)^2}{\sigma^2} \right)$$

- ▶ We want to estimate W from n *iid* observations $\{y_i(X_i), i = 1, \dots, n\}$.

- We take the prior density of W as

$$f(W) = \left(\frac{1}{\alpha\sqrt{2\pi}} \right)^d \exp \left(-\frac{W^T W}{2\alpha^2} \right)$$

which is zero-mean normal with diagonal covariance matrix; α is a hyper-parameter.

Now the posterior density is given by

$$\begin{aligned} f(W | Y) &\propto \prod_{i=1}^n f(y_i | X_i, W, \sigma) f(W) \\ &\propto \exp \left(- \sum_{i=1}^n \frac{(y_i - W^T X_i)^2}{2\sigma^2} - \frac{1}{2\alpha^2} W^T W \right) \end{aligned}$$

- ▶ To find the MAP estimate we need to maximize the posterior density
- ▶ We can maximize log of the posterior.

The log posterior is of the form

$$\ln(f(W | Y)) = -\frac{1}{2} \sum_{i=1}^n (y_i - W^T X_i)^2 - \lambda W^T W + K$$

- ▶ Maximizing this is same as minimizing the regularized least squares criterion.
- ▶ Hence the MAP estimate is the regularized least squares solution.

Bayesian Linear Regression

- ▶ We saw that we can look at linear least squares estimation in a Bayesian framework also.
- ▶ The MAP estimate corresponds to regularized least squares.
- ▶ MAP is a point estimate obtained from the posterior.
- ▶ In Bayesian framework, the estimate is whole of the posterior.
- ▶ We can also calculate the distribution of prediction variable based on data.

The posterior density

- ▶ In general, the prior can be

$$f(W) = \mathcal{N}(W | \boldsymbol{\mu}_0, \Sigma_0)$$

- ▶ The probability model for y is

$$f(Y|A, W, \sigma^2) = \prod_{i=1}^n \mathcal{N}(y_i | X_i^T W, \sigma^2) = \mathcal{N}(Y | AW, \sigma^2 I)$$

- ▶ The posterior would also be a Gaussian which can be calculated using the technique of ‘completing the squares’

The posterior density

- ▶ The posterior would be

$$f(W|Y, A, \sigma^2) = \mathcal{N}(W|\boldsymbol{\mu}_n, \Sigma_n)$$

where

$$\boldsymbol{\mu}_n = \Sigma_n (\sigma^2 A^T Y + \Sigma_0^{-1} \boldsymbol{\mu}_0)$$

$$\Sigma_n^{-1} = (\Sigma_0^{-1} + \sigma^2 A^T A)$$

The Predictive Distribution

- ▶ We can also Calculate (for any given new X)

$$f(y|X, A, Y, \sigma^2) = \int f(y|X, W, \sigma^2) f(W|A, Y, \sigma^2) dW$$

which is $\mathcal{N}(y|\boldsymbol{\mu}_n^T X, A^T \Sigma_n A + \sigma^2)$.

- ▶ We can use this to predict the y for any X .

- ▶ For linear regression we want to minimize

$$J(W) = (1/2) \sum_i (W^T X_i - y_i)^2$$

- ▶ We can think of this as

$$J(W) = \sum_i L(W^T X_i, y_i)$$

where L is a loss function.

- ▶ The squared error is only one choice for the loss.
- ▶ We can have others, e.g., absolute value of error.

Robust Regression

- ▶ The squared error criterion is sensitive to outliers. The absolute value of error is more robust to outliers.
- ▶ There are multiple viewpoints possible on this.
- ▶ One way of understanding it is that mean is more sensitive to outliers compared to median.
- ▶ Another way of looking at it is that the Gaussian noise is a 'light-tailed' distribution. (Recall that we get the criterion of squared error when we assume Gaussian noise and estimate a discriminative model).

Robust Regression

- ▶ The squared error criterion is sensitive to outliers. The absolute value of error is more robust to outliers.
- ▶ There are multiple viewpoints possible on this.
- ▶ One way of understanding it is that mean is more sensitive to outliers compared to median.
- ▶ Another way of looking at it is that the Gaussian noise is a 'light-tailed' distribution.
- ▶ Instead of Gaussian noise we can assume Laplace distribution for noise:

$$f_{\text{Lap}}(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right), \quad -\infty < x < \infty$$

- ▶ Now maximizing likelihood would result in minimizing absolute value of error.
- ▶ This is a 'heavy-tailed' distribution.

- ▶ Suppose we want to learn W to minimize

$$J(W) = \sum_{i=1}^n | (y_i - W^T X_i) | = \| (AW - Y) \|_1$$

- ▶ Optimizing absolute value of error is difficult.
- ▶ The objective function is convex but non-smooth.
- ▶ We can use some convex optimization techniques.
- ▶ Another interesting method for this is Iterated Reweighted Least Squares (IRLS).

Weighted Least Squares

- ▶ Suppose we want to minimize

$$J(W) = \sum_{i=1}^n b_i (W^T X_i - y_i)^2, \quad b_i > 0$$

- ▶ Define a diagonal matrix $B = \text{diag}(\sqrt{b_i})$.
- ▶ Now minimizing J is same as least squares where we multiply X_i and y_i by B_{ii} .
- ▶ With A and Y as earlier, let $\tilde{A} = BA$ and $\tilde{Y} = BY$.
- ▶ So, now we get the solution as $W^* = (\tilde{A}^T \tilde{A})^{-1} \tilde{A}^T \tilde{Y}$.
- ▶ Thus, we can easily solve the least squares problem with any positive weights.

- ▶ Suppose we know how to minimize over W (for any b)

$$C(W, b) = \sum_{i=1}^n b_i g_i(W) \quad [g_i(W) = (W^T X_i - y_i)^2]$$

- ▶ We actually want to minimize

$$C_h(W) = \sum_{i=1}^n h(g_i(W)) \quad [h(x) = \sqrt{x}]$$

- ▶ Heuristically, we want to ensure descent on C should be descent on C_h .

$$\nabla_W C(W, b) = \nabla_W C_h(W) \Rightarrow b_i \nabla_W g_i(W) = h'(g_i(W)) \nabla_W g_i(W)$$

- ▶ Hence, we can take $b_i = h'(g_i(W))$.
(In the iterative method, we use previous iteration value of W).

IRLS algorithm for minimizing absolute error

- ▶ We want to minimize

$$J(W) = \sum_{i=1}^n h(g_i(W)) = \sum_{i=1}^n \sqrt{(W^T X_i - y_i)^2}$$

- ▶ Note that for $h(x) = \sqrt{x}$, $h'(x) = 0.5(x)^{-0.5}$
- ▶ Hence, we get the iterative method as

$$W(k+1) = \arg \min_W \sum_{i=1}^n b_i^k (W^T X_i - y_i)^2$$

where

$$b_i^k = h'(g_i(W(k))) = \frac{1}{2} |W(k)^T X_i - y_i|^{-1}$$

- ▶ We solve the minimization using least squares.
- ▶ One can show this is a descent method for J .

- ▶ Our notation: $g_i(W) = (W^T X_i - y_i)^2$
- ▶ We want to minimize $J(W) = \sum_{i=1}^n h(g_i(W))$ where h is concave. (For us $h(x) = \sqrt{x}$).
- ▶ Let us write: $r_i^k = g_i(W(k))$, $r_i^{k+1} = g_i(W(k+1))$
- ▶ Our iterations are

$$W(k+1) = \arg \min_W \sum_{i=1}^n h'(r_i^k) g_i(W)$$

- ▶ Hence we have

$$\sum_{i=1}^n h'(r_i^k) r_i^{k+1} < \sum_{i=1}^n h'(r_i^k) r_i^k$$

- ▶ Using this we need to show

$$J(W(k+1)) = \sum_{i=1}^n h(r_i^{k+1}) < \sum_{i=1}^n h(r_i^k) = J(W(k))$$

- ▶ Since h is concave, we have

$$h(r_i^{k+1}) \leq h(r_i^k) + h'(r_i^k)(r_i^{k+1} - r_i^k)$$

- ▶ By summing the above for $i = 1$ to n

$$\sum_{i=1}^n h(r_i^{k+1}) \leq \sum_{i=1}^n h(r_i^k) + K$$

where

$$K = \sum_{i=1}^n h'(r_i^k) r_i^{k+1} - \sum_{i=1}^n h'(r_i^k) r_i^k < 0$$

- ▶ Hence we get

$$J(W(k+1)) = \sum_{i=1}^n h(r_i^{k+1}) < \sum_{i=1}^n h(r_i^k) = J(W(k))$$

Huber Loss

- ▶ The squared error criterion is easy to optimize but is sensitive to outliers.
- ▶ The absolute value of error is more robust to outliers. But the optimization problem is more difficult because the absolute value function is non-differentiable at zero.
- ▶ An interesting approach is to use a hybrid of the two – use Huber loss defined by

$$\begin{aligned} L_{H_\delta}(a, b) &= 0.5(a - b)^2 && \text{if } |a - b| \leq \delta \\ &= \delta|a - b| - 0.5\delta^2 && \text{if } |a - b| > \delta \end{aligned}$$

- ▶ We learn W to minimize $J(W) = \sum_i L_{H_\delta}(y_i, W^T X_i)$.
- ▶ Using $\frac{d}{dr}|r| = \text{sign}(r)$, this loss function is differentiable and hence the resulting optimization problem has a smooth objective function.

Fisher Linear Discriminant

- ▶ Minimizing mean squared error is one of the possible methods for learning linear classifiers and regressors.
- ▶ Fisher Linear Discriminant is another way of constructing linear classifiers.


Fisher Linear Discriminant

- ▶ A linear discriminant function based classifier is:


$$\text{Decide } X \in C-1 \text{ if } W^T X + w_0 > 0$$

- ▶ We project the data along the direction W .
- ▶ Hence One can think of the best W as the direction along which the two classes are well separated.
- ▶ Such a method is called Fisher Linear Discriminant.

- Consider the following 2-class example



- ▶ A good direction to project the data here is as shown.



- ▶ Fisher Linear Discriminant is based on formalizing this notion

Fisher Linear Discriminant

- ▶ The idea is to find a direction W such that the training data of the two classes are well-separated if projected onto this direction.
- ▶ We consider the 2-class case.

- ▶ Let $\{(X_i, y_i), i = 1, \dots, n\}$ be the data.
- ▶ Let $y_i \in \{0, 1\}$.
- ▶ Let C_0 and C_1 denote the two classes. Thus, if $y_i = 0$ then $X_i \in C_0$ and if $y_i = 1$ then $X_i \in C_1$.
- ▶ Let n_0 and n_1 denote the number of examples of each class. ($n = n_0 + n_1$)
- ▶ For any W , let $z_i = W^T X_i$.
- ▶ z_i are the one dimensional data that we get after projection.

- ▶ Let M_0 and M_1 be the means of data from the two classes:

$$M_0 = \frac{1}{n_0} \sum_{X_i \in C_0} X_i; \quad M_1 = \frac{1}{n_1} \sum_{X_i \in C_1} X_i$$

The corresponding means of the projected data would be

$$m_0 = W^T M_0 \quad \text{and} \quad m_1 = W^T M_1$$

- ▶ The difference $(m_0 - m_1)$ gives us an idea of the separation between samples of the two classes after projecting the data onto the direction W .
- ▶ Hence, we may want a W that maximizes $(m_0 - m_1)^2$.
- ▶ However, we have to make this scale independent.
- ▶ Also, the distance between means should be viewed relative to the variances.

- Define

$$s_0^2 = \sum_{X_i \in C_0} (W^T X_i - m_0)^2; \quad s_1^2 = \sum_{X_i \in C_1} (W^T X_i - m_1)^2$$

These give us the variances (upto a factor) of the two classes in the projected data.

- We want large separation between m_0 and m_1 relative to the variances.

- ▶ Hence we can take our objective to be to maximize

$$J(W) = \frac{(m_1 - m_0)^2}{s_0^2 + s_1^2}$$

- ▶ We now rewrite J into a more convenient form.
- ▶ We have

$$\begin{aligned}(m_1 - m_0)^2 &= (W^T M_1 - W^T M_0)^2 \\ &= [W^T (M_1 - M_0)][W^T (M_1 - M_0)]^T \\ &= W^T (M_1 - M_0)(M_1 - M_0)^T W\end{aligned}$$

- ▶ This is a popular 'trick'. In general
 $(X^T Y)^2 = (X^T Y)(Y^T X) = X^T Y Y^T X$

- ▶ Thus we have $(m_1 - m_0)^2 = W^T S_B W$ where

$$S_B = (M_1 - M_0)(M_1 - M_0)^T$$

is a $d \times d$ matrix (note that $X_i \in \Re^d$).

- ▶ It is called *between class* scatter matrix.
- ▶ We can similarly write s_0^2 and s_1^2 also as quadratic forms.

We have

$$\begin{aligned}s_0^2 &= \sum_{X_i \in C_0} (W^T X_i - W^T M_0)^2 \\&= \sum_{X_i \in C_0} [W^T (X_i - M_0)]^2 \\&= \sum_{X_i \in C_0} W^T (X_i - M_0) (X_i - M_0)^T W \\&= W^T \left[\sum_{X_i \in C_0} (X_i - M_0) (X_i - M_0)^T \right] W\end{aligned}$$

- ▶ Similarly, we get

$$s_1^2 = W^T \left[\sum_{X_i \in C_1} (X_i - M_1)(X_i - M_1)^T \right] W$$

- ▶ Thus we can write $s_0^2 + s_1^2 = W^T S_w W$, where

$$S_w = \sum_{X_i \in C_0} (X_i - M_0)(X_i - M_0)^T + \sum_{X_i \in C_1} (X_i - M_1)(X_i - M_1)^T$$

- ▶ S_w is also $d \times d$ matrix and is called *within class* scatter matrix.

- ▶ Hence we can now write J as

$$J(W) = \frac{W^T S_B W}{W^T S_w W}$$

- ▶ We want to find a W that maximizes $J(W)$.
- ▶ Note that $J(W)$ is not affected by scaling of W .
- ▶ Given the data we can calculate the S_B and S_w .
- ▶ Maximizing ratio of quadratic forms is a standard optimization problem.

- ▶ We need to maximize

$$J(W) = \frac{W^T S_B W}{W^T S_w W}$$

- ▶ Differentiating w.r.t. W and equating to zero, we get

$$\frac{2S_B W}{W^T S_w W} - \frac{W^T S_B W}{(W^T S_w W)^2} 2S_w W = 0$$

- ▶ Implies, $S_B W$ is in the same direction as $S_w W$.

- ▶ Thus, any maximizer of $J(W)$ has to satisfy

$$S_w W = \lambda S_B W$$

for some constant λ .

- ▶ This is known as the generalized eigen value problem.
- ▶ There are standard methods to solve this problem using, e.g., LU decomposition.
- ▶ By solving the generalized eigen value problem we can find the best direction W .

- ▶ Often, the real symmetric matrix S_w would be invertible.
- ▶ Recall that

$$S_w = \sum_{X_i \in C_0} (X_i - M_0)(X_i - M_0)^T + \sum_{X_i \in C_1} (X_i - M_1)(X_i - M_1)^T$$

- ▶ This is a sum of large number of rank 1 matrices.
- ▶ Also each term in S_w is proportional to the sample-mean-estimate of the covariance matrix of one of the class conditional densities.

- ▶ If S_w is invertible, then we can write

$$W = S_w^{-1} S_B W$$

- ▶ We have

$$S_B W = (M_1 - M_0)(M_1 - M_0)^T W = k(M_1 - M_0)$$

where k is some constant. (note $k = (m_1 - m_0)$)

- ▶ Now we get (since scale factor in W is not relevant)

$$W = S_w^{-1}(M_1 - M_0)$$

Obtaining Fisher Linear Discriminant

- ▶ We can sum-up the process as follows.
- ▶ Given the training data, we first form the scatter matrix S_w and also calculate the means M_0 and M_1 .
- ▶ If S_w is invertible, we calculate W by $W = S_w^{-1}(M_1 - M_0)$.
- ▶ Even if S_w is not invertible, there are techniques to find the maximizer of $J(W)$ by solving the generalized eigen value problem.
- ▶ Thus we can find the best direction W .

- ▶ The final (linear) classifier is $\text{sign}(W^T X + b)$.
- ▶ So far, we have seen how to obtain best W .
- ▶ We still have to learn the best b also. But this is a simple threshold learning problem.
- ▶ We can do a simple line search to find the threshold b to maximize probability of correct classification.
- ▶ Or, we can take the one dimensional projected data and learn the best classifier by, e.g., modelling the class conditional densities as normal.

- ▶ Fisher Linear Discriminant is also a popular classifier.
- ▶ Though the method looks quite different from that of linear least squares there are close connections between the two.

- ▶ Given the original training data $\{(X_i, y_i)\}$ we form new training data $\{(X_i, y'_i)\}$ as follows.
- ▶ We take $y'_i = n/n_0$ if $y_i = 0$ and $y'_i = -n/n_1$ if $y_i = 1$.
- ▶ We now treat this as a data for a regression problem and learn a model $\hat{y} = W^T X + b$ using linear least squares.
- ▶ It can be shown that the least squares solution for W would be same as that of FLD.
- ▶ Thus FLD can be viewed as a special case of linear least squares.

- ▶ We have considered various methods of learning linear classifiers and regression functions.
- ▶ In the case of regression, we have considered only estimating real-valued functions. We can generalize this to vector-valued functions.
- ▶ In classification we considered only 2-class problems. This can also be generalized to multi-class case.


- ▶ First consider estimating vector-valued functions.
- ▶ Now the training data is $\{(X_i, y_i), i = 1, \dots, n\}$ where $X_i \in \mathbb{R}^d$ and $y_i = (y_i^1, \dots, y_i^m) \in \mathbb{R}^m$.
- ▶ For any given X we want to predict the target $y = (y^1, \dots, y^m)$.
- ▶ Thus we want to learn $W_j, b_j, j = 1, \dots, m$ so that

$$\hat{y}^j = W_j^T X + b_j, j = 1, \dots, m$$


- ▶ We can obtain these by simply solving m number of linear least squares regression problems.

- ▶ Now let us consider the multi-class problem.
- ▶ Suppose we have K classes: C_1, \dots, C_K .
- ▶ We can solve the multi-class problem by learning a number of 2-class classifiers.
- ▶ For example, we can learn K two class classifiers: ' C_i Vs not- C_i '. (Called one versus rest).
- ▶ Or we can learn $K(K - 1)/2$ number of 2-class classifiers: ' C_i Vs C_j '
- ▶ But neither of these approaches are really satisfactory for generalizing linear discriminant functions.

- Suppose we try 'one versus rest' approach. Then there may be regions of feature space where classification is ambiguous.



- ▶ Similar problem exists when we try ' C_i Vs C_j ' approach



- ▶ A better way to formulate a linear discriminant function based classifier for the multi-class case is as follows.
- ▶ We will have K functions, g_s , $s = 1, \dots, K$, given by

$$g_s(X) = W_s^T X + b_s$$

- ▶ Now the classifier would assign class C_j to X if

$$g_j(X) \geq g_s(X), \forall s$$

- ▶ Essentially, we are approximating the posterior probability of j^{th} class by g_j .
- ▶ Recall that this is the structure of Bayes classifier in multi-class case.
- ▶ In the above we would have a fixed (may be arbitrary) rule for breaking ties.

- ▶ Now, to learn a linear classifier for the K -class case, we need to learn the K functions g_s .
- ▶ The simplest way to do this is to make the class label to be a vector of K components.
- ▶ If $X_i \in C_j$ then y_i would be a K -vector with j^{th} component one and all others zero.
- ▶ Now learning the K functions is same as linear regression with vector valued targets.

Logistic Regression – multi-class case

- ▶ We can similarly generalize logistic regression also for multi-class case.
- ▶ Let us recall the main idea in logistic regression in the two class case.
- ▶ We approximate posterior probability as

$$q_1(X) = h(W^T X + w_0)$$

where h is the logistic function

$$h(a) = \frac{1}{1 + \exp(-a)}$$

- The motivation for using the logistic function is

$$\begin{aligned} q_1(X) &= \frac{f_1(X) p_1}{f_0(X) p_0 + f_1(X) p_1} \\ &= \frac{1}{1 + \exp(-\xi)} \quad \text{where} \end{aligned}$$

$$\xi = - \ln \left(\frac{f_0(X) p_0}{f_1(X) p_1} \right) = \ln \left(\frac{f_1(X) p_1}{f_0(X) p_0} \right)$$

- We now use the same Bayes rule to find a convenient model for posterior probabilities in the multiclass case.

- ▶ In the multi-class case, Bayes rule gives

$$q_j(X) = \frac{f_j(X)p_j}{\sum_s f_s(X)p_s} = \frac{\exp(a_j)}{\sum_s \exp(a_s)}$$

where $a_s = \ln(f_s(X)p_s)$.

- ▶ The idea is that we approximate $a_s = W_s^T X + w_{s0}$.
- ▶ The above function is a good candidate for modeling posterior probabilities in the multi-class case.

- ▶ In the two class case we want to know which of $f_1(X)p_1$ and $f_0(X)p_0$ is greater.
- ▶ This can be done by looking at sign of $\ln \left(\frac{f_1(X)p_1}{f_0(X)p_0} \right)$.
- ▶ 'sign' is a discontinuous function and the logistic function is a kind of continuous analog for this.
- ▶ In the multiclass case, we need to find the maximum of $f_i(X)p_i$, $i = 1, \dots, K$.
- ▶ So, we need a smooth function to approximate the maximum computation.

- ▶ Define a function $g : \Re^K \rightarrow \Re^K$, with $g(a) = [g_1(a) \cdots g_K(a)]^T$ and for $j = 1, \dots, K$,

$$g_j(a) = \frac{\exp(a_j)}{\sum_s \exp(a_s)}, \quad a = (a_1, \dots, a_K)^T \in \Re^K.$$

- ▶ This is known as the softmax function.
- ▶ Essentially if a_j is the maximum of the components of a then $g_j(a)$ would be closer to one and all other components of g would be closer to zero.
- ▶ We note, for later use, that

$$\frac{\partial g_k}{\partial a_j} = g_k(a)(\delta_{kj} - g_j(a))$$

where $\delta_{kj} = 0$ if $k \neq j$ and $\delta_{kj} = 1$ if $k = j$.

- ▶ We now take, for each s , $a_s = W_s^T X + w_{s0}$ and learn all W_s and w_{s0} .
- ▶ Using augmented feature vector, we can write $a_s = W_s^T X$. Let W be a matrix with columns W_s .
- ▶ After learning, W , given a new X , we calculate $g(W^T X)$ and then put X in class C_j if the j^{th} component of $g(W^T X)$ is the highest.

- ▶ In the 2-class logistic regression, we used the logistic function to model posterior probability.
- ▶ In the multi-class case we want to use the softmax function for modeling the conditional distribution of y given X .
- ▶ Let us now write

$$g_j(W, X) = \frac{\exp(W_j^T X)}{\sum_s \exp(W_s^T X)}$$

- ▶ Let us represent the class label as a one-hot vector:

$$y = (y^1, \dots, y^K)^T; \quad y^j \in \{0, 1\}; \quad \sum_j y^j = 1$$

- ▶ y takes only K different values and let us represent them as e_1, \dots, e_K .
- ▶ We take (as our probability model)
 $P[y = e_j \mid X] = g_j(W, X)$.

- ▶ The probability model for the conditional distribution now is

$$f(y \mid X, W) = \prod_{i=1}^K (g_i(W, X))^{y_i}$$

- ▶ Let the data be $\mathcal{D} = \{(X_1, y_1), \dots, (X_n, y_n)\}$.
- ▶ The likelihood now is

$$L(W_1, \dots, W_K \mid \mathcal{D}) = \prod_{i=1}^n f(y_i \mid X_i, W) = \prod_{i=1}^n \prod_{j=1}^K (g_j(W, X_i))^{y_i^j}$$

- ▶ The log likelihood is

$$l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n \sum_{j=1}^K y_i^j \ln (g_j(W, X_i))$$

- ▶ We need to maximize this to learn the W .

- ▶ The log likelihood is

$$l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n \sum_{j=1}^K y_i^j \ln (g_j(W, X_i))$$

- ▶ By differentiating this and using earlier formula for derivative of g_j , we can show that

$$\nabla_{W_j} l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n (y_i^j - g_j(W, X_i)) X_i$$

- ▶ Hence an iterative algorithm for ML estimate of W_j , $j = 1, \dots, K$, is

$$W_j(k+1) = W_j(k) + \eta \sum_{i=1}^n (y_i^j - g_j(W, X_i)) X_i$$

- ▶ This is the multi-class logistic regression.

- ▶ We can generalize Fisher linear discriminant also to multi-class case.
- ▶ In the 2-class case we are interested in finding a direction or a one-dimensional subspace onto which we project the data.
- ▶ In the K -class case, we want to find a $(K - 1)$ -dimensional subspace onto which we project the data.
- ▶ The idea is to find a subspace where in the projected data, the means of the two classes have maximum separation relative to the variances.

Learning and generalization

- ▶ The problem of designing a classifier is essentially one of learning from examples.
- ▶ Given training data, we want to find an appropriate classifier.
- ▶ It amounts to searching over a family of classifiers to find one that minimizes 'error' over training set.
- ▶ For example, in least squares approach we are searching over the family of linear classifiers for minimizing square of error.

- ▶ As we discussed earlier, performance on training set is not the real issue.
- ▶ We would like the learnt classifier to perform well on new data.
- ▶ This is the issue of generalization. Does the learnt classifier generalize well?

- ▶ In practice one assesses the generalization of the learnt classifier by looking at the error on a separate set of labelled data called test set.
- ▶ Since the test set would not be used in training, error on that data could be a good measure of the performance of the learnt classifier.
- ▶ But here we are more interested in formalizing the notion of generalization error.
- ▶ we look at the specific issues of practice later on.
Currently our focus would be on theoretical analysis of how to say whether a learning algorithm would generalize well.

- ▶ We can see the main issue through a simple example of regression
- ▶ Suppose we have data $\{(X_i, y_i)\}$, $X_i, y_i \in \mathbb{R}$.
- ▶ We want to learn a function f so that we can predict y as $f(X)$.
- ▶ This is a simple regression problem and we can use least squares for it based on the form of f .

- ▶ Suppose we choose polynomial function

$$f(X) = w_0 + w_1X + w_2X^2 + \cdots + w_mX^m$$

- ▶ As we discussed earlier we can easily learn this using linear least squares algorithm.
- ▶ One question is what m to choose.
- ▶ We have looked at regularized least squares for this. (It does not tell best m but helps learn a model with good generalization for a given m).
- ▶ There are other methods (e.g., BIC)

- ▶ But specifically, let us ask can our data error tell what m is proper.
- ▶ Firstly the fact that we get less error for m' compared to m does not necessarily mean m' -degree is a better fit.
- ▶ For a particular m if we get very low data error, can we say it is good?
- ▶ We know that if we search over all polynomials, we can never really learn anything. Can we formalize such notions precisely?
- ▶ There are different ways of addressing this issue (MDL, VC-theory etc).

- ▶ Any learning algorithm takes training data as the input and outputs a specific classifier/function.
- ▶ For this, it searches over some chosen family of functions to find one that optimizes a chosen criterion function.

$$\{(X_i, y_i)\} \rightarrow \boxed{\begin{array}{c} \text{Learning Algorithm} \\ \text{(searching over } \mathcal{F}) \end{array}} \rightarrow f \in \mathcal{F}$$

- ▶ The question is: how can we formalize correctness of learning?
- ▶ There are many ways of addressing this issue (MDL, VC-theory etc).

- ▶ For example, a generic approach is what is called **Minimum Description Length** principle.
- ▶ Suppose we want to send the data over a communication channel.
- ▶ we can send the $2n$ numbers, X_i, y_i using some number of bits.
- ▶ Or we can send X_i , the function f and the errors $y_i - f(X_i)$.

- ▶ If the fit is good, the errors $y_i - f(X_i)$ would be small and we may be able to send them using smaller number of bits compared sending y_i .
- ▶ However, we also need to send f .
- ▶ If f is very complex, then what we save in bits by sending errors instead of y_i may be more than offset by the bits needed to send description of f .
- ▶ Hence we can rate different f by the total number of bits we need.
- ▶ This can balance the data error and model complexity in a natural way.

- ▶ We will follow a different statistical approach to address the issue of correctness of learning.
- ▶ We begin with a simple formalism where there is no 'noise' and the goal of learning is well-defined.