

Summary – Linear Classification / Regression

- ▶ A linear classifier is given by:

$$\begin{aligned} h(X) &= 1 \text{ if } \sum_{i=1}^{d'} w_i \phi_i(X) + w_0 > 0 \\ &= 0 \text{ Otherwise} \end{aligned}$$

where ϕ_i are fixed functions.

- ▶ For simplicity of notation, we use augmented feature vector and omit ϕ .
- ▶ We take $h(X) = \text{sign}(W^T X)$ for simplicity of notation.
- ▶ Perceptron is a classical algorithm to learn such a classifier.

Summary – Linear Classification / Regression

- ▶ For linear regression the objective is to learn a model:

$$\hat{y}(X) = \sum_{i=1}^d w_i x_i + w_0 = W^T X$$

We could use $\phi_i(X)$ in place of x_i .

- ▶ The least squares criterion is to minimize

$$J(W) = \frac{1}{2} \sum_{i=1}^n (W^T X_i - y_i)^2$$

- ▶ The minimizer of J is given by

$$W^* = (A^T A)^{-1} A^T Y$$

Where A is a matrix whose rows are X_i^T and Y is a vector whose components are y_i .

- ▶ This method can be used for learning linear classifiers also

Summary – Linear Classification / Regression

- ▶ We can also minimize J by iterative gradient descent.
- ▶ An incremental version of this gradient descent is the LMS algorithm.
- ▶ The LMS algorithm is:

$$W(k+1) = W(k) - \eta (X(k)^T W(k) - y(k)) X(k)$$

where $(X(k), y(k))$ is the (random) sample picked and $W(k)$ is the weight vector at iteration k .

- ▶ This is used in many adaptive signal processing problems.

Summary – Linear Classification / Regression

- ▶ For any random variables X, y

$$f^*(X) \stackrel{\text{def}}{=} \arg \min_f E[(y - f(X))^2] = E[y|X]$$

- ▶ If X, y are jointly Gaussian then this is a linear function.
- ▶ If $y \in \{0, 1\}$, then

$$E[y|X] = \text{Prob}[y = 1|X] = q_1(X)$$

- ▶ In linear least squares method, we are essentially learning a parametric model for the posterior probability (A discriminative model)

Summary – Linear Classification / Regression

- ▶ If we assume the model $y = W^T X + \xi$ where ξ is a zero-mean Gaussian, then the MLE would be same as the linear least squares solution.
- ▶ The residual squared error is the MLE of noise variance.
- ▶ For linear regression, the Gaussian noise assumption is reasonable.

Summary – Linear Classification / Regression

- ▶ For $y \in \{0, 1\}$, we can model conditional distribution of y given X as Bernoulli with parameter $\sigma(W^T X)$.
- ▶ The resulting algorithm to maximize log likelihood is logistic regression:

$$W(k+1) = W(k) + \eta \sum_{i=1}^n (y_i - \sigma(W^T X_i)) X_i$$

- ▶ Logistic regression is a better way to learn posterior probability for classification problems.
- ▶ Logistic regression is a prototypical example of learning a discriminative model.
- ▶ Logistic regression can also be formulated as an IRLS algorithm.

Summary – Linear Classification / Regression

- ▶ In regularized least squares (Ridge regression) we minimize

$$J(W) = \frac{1}{2} \sum_{i=1}^n (W^T X_i - y_i)^2 + \frac{\lambda}{2} W^T W$$

- ▶ The solution now is

$$W^* = (A^T A + \lambda I)^{-1} A^T Y$$

- ▶ In general, in regularization we take the objective to be sum of data error and model complexity.
- ▶ Regularization helps mitigate problems of overfitting.

Summary – Linear Classification / Regression

- ▶ Under the model $y = W^T X + \xi$ where ξ is a zero-mean Gaussian, the MLE is the linear least squares solution.
- ▶ Under the same model, the MAP estimate with a Gaussian prior is the L_2 -regularized least squares solution.
- ▶ In general, different regularization terms amount to choosing different priors.

Summary – Linear Classification / Regression

- ▶ Least squares soln is sensitive to outliers.
- ▶ Minimizing absolute value of error is more robust.
- ▶ We can formulate it as an IRLS algorithm.
- ▶ Huber loss is combination of squared and absolute value of error and is useful for robust regression.

Summary – Linear Classification / Regression

- ▶ Fisher Linear Discriminant is another way to learn a linear classifier.
- ▶ Seeks to find a direction along which the projected data has best separation between the two classes.
- ▶ It can be obtained by solving a generalized eigen value problem

Summary – Linear Classification / Regression

- ▶ Generalizing linear regression to handle vector-valued functions is straight-forward.
- ▶ We can formulate a K -class linear classifier by having K functions $g_s(X) = W_s^T X + b_s$ and having

$$h(X) = C_j \text{ if } g_j(X) \geq g_s(X), \forall s$$

- ▶ Then learning multi-class classifier through least squares method is same as learning vector-valued functions.

Logistic Regression – multi-class case

- ▶ We can similarly generalize logistic regression also for multi-class case.
- ▶ Let us recall the main idea in logistic regression in the two class case.
- ▶ We approximate posterior probability as

$$q_1(X) = h(W^T X + w_0)$$

where h is the logistic function

$$h(a) = \frac{1}{1 + \exp(-a)}$$

- The motivation for using the logistic function is

$$\begin{aligned} q_1(X) &= \frac{f_1(X) p_1}{f_0(X) p_0 + f_1(X) p_1} \\ &= \frac{1}{1 + \exp(-\xi)} \quad \text{where} \end{aligned}$$

$$\xi = -\ln \left(\frac{f_0(X) p_0}{f_1(X) p_1} \right) = \ln \left(\frac{f_1(X) p_1}{f_0(X) p_0} \right)$$

- We now use the same Bayes rule to find a convenient model for posterior probabilities in the multiclass case.

- ▶ In the multi-class case, Bayes rule gives

$$q_j(X) = \frac{f_j(X)p_j}{\sum_s f_s(X)p_s} = \frac{\exp(a_j)}{\sum_s \exp(a_s)}$$

where $a_s = \ln(f_s(X)p_s)$.

- ▶ The idea is that we approximate $a_s = W_s^T X + w_{s0}$.
- ▶ The above function is a good candidate for modeling posterior probabilities in the multi-class case.

- ▶ In the two class case we want to know which of $f_1(X)p_1$ and $f_0(X)p_0$ is greater.
- ▶ This can be done by looking at sign of $\ln \left(\frac{f_1(X)p_1}{f_0(X)p_0} \right)$.
- ▶ 'sign' is a discontinuous function and the logistic function is a kind of continuous analog for this.
- ▶ In the multiclass case, we need to find the maximum of $f_i(X)p_i$, $i = 1, \dots, K$.
- ▶ So, we need a smooth function to approximate the maximum computation.

- ▶ Define a function $g : \Re^K \rightarrow \Re^K$, with $g(a) = [g_1(a) \cdots g_K(a)]^T$ and for $j = 1, \dots, K$,

$$g_j(a) = \frac{\exp(a_j)}{\sum_s \exp(a_s)}, \quad a = (a_1, \dots, a_K)^T \in \Re^K.$$

- ▶ This is known as the softmax function.
- ▶ Essentially if a_j is the maximum of the components of a then $g_j(a)$ would be closer to one and all other components of g would be closer to zero.
- ▶ We note, for later use, that

$$\frac{\partial g_k}{\partial a_j} = g_k(a)(\delta_{kj} - g_j(a))$$

where $\delta_{kj} = 0$ if $k \neq j$ and $\delta_{kj} = 1$ if $k = j$.

- ▶ We now take, for each s , $a_s = W_s^T X + w_{s0}$ and learn all W_s and w_{s0} .
- ▶ Using augmented feature vector, we can write $a_s = W_s^T X$. Let W be a matrix with columns W_s .
- ▶ After learning, W , given a new X , we calculate $g(W^T X)$ and then put X in class C_j if the j^{th} component of $g(W^T X)$ is the highest.

- ▶ In the 2-class logistic regression, we used the logistic function to model posterior probability.
- ▶ In the multi-class case we want to use the softmax function for modeling the conditional distribution of y given X .
- ▶ Let us now write

$$g_j(W, X) = \frac{\exp(W_j^T X)}{\sum_s \exp(W_s^T X)}$$

- ▶ Let us represent the class label as a one-hot vector:

$$y = (y^1, \dots, y^K)^T; \quad y^j \in \{0, 1\}; \quad \sum_j y^j = 1$$

- ▶ y takes only K different values and let us represent them as e_1, \dots, e_K .
- ▶ We take (as our probability model)
 $P[y = e_j \mid X] = g_j(W, X)$.

- ▶ The probability model for the conditional distribution now is

$$f(y \mid X, W) = \prod_{i=1}^K (g_i(W, X))^{y_i}$$

- ▶ Let the data be $\mathcal{D} = \{(X_1, y_1), \dots, (X_n, y_n)\}$.
- ▶ The likelihood now is

$$L(W_1, \dots, W_K \mid \mathcal{D}) = \prod_{i=1}^n f(y_i \mid X_i, W) = \prod_{i=1}^n \prod_{j=1}^K (g_j(W, X_i))^{y_i^j}$$

- ▶ The log likelihood is

$$l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n \sum_{j=1}^K y_i^j \ln (g_j(W, X_i))$$

- ▶ We need to maximize this to learn the W .

- ▶ The log likelihood is

$$l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n \sum_{j=1}^K y_i^j \ln (g_j(W, X_i))$$

- ▶ By differentiating this and using earlier formula for derivative of g_j , we can show that

$$\nabla_{W_j} l(W_1, \dots, W_K \mid \mathcal{D}) = \sum_{i=1}^n (y_i^j - g_j(W, X_i)) X_i$$

- ▶ Hence an iterative algorithm for ML estimate of W_j , $j = 1, \dots, K$, is

$$W_j(k+1) = W_j(k) + \eta \sum_{i=1}^n (y_i^j - g_j(W, X_i)) X_i$$

- ▶ This is the multi-class logistic regression.

- ▶ We can generalize Fisher linear discriminant also to multi-class case.
- ▶ In the 2-class case we are interested in finding a direction or a one-dimensional subspace onto which we project the data.
- ▶ In the K -class case, we want to find a $(K - 1)$ -dimensional subspace onto which we project the data.
- ▶ The idea is to find a subspace where in the projected data, the means of the two classes have maximum separation relative to the variances.

Learning and generalization

- ▶ The problem of designing a classifier is essentially one of learning from examples.
- ▶ Given training data, we want to find an appropriate classifier.
- ▶ It amounts to searching over a family of classifiers to find one that minimizes 'error' over training set.
- ▶ For example, in least squares approach we are searching over the family of linear classifiers for minimizing square of error.

- ▶ As we discussed earlier, performance on training set is not the ultimate objective.
- ▶ We would like the learnt classifier to perform well on new data.
- ▶ This is the issue of generalization. Does the learnt classifier generalize well?

- ▶ In practice one assesses the generalization of the learnt classifier by looking at the error on a separate set of labelled data called test set.
- ▶ Since the test set would not be used in training, error on that data could be a good measure of the performance of the learnt classifier.
- ▶ But here we are more interested in formalizing the notion of generalization error.
- ▶ We look at the specific issues of practice later on.
Currently our focus would be on theoretical analysis of how to say whether a learning algorithm would generalize well.

- ▶ Any learning algorithm takes training data as the input and outputs a specific classifier/function.
- ▶ For this, it searches over some chosen family of functions to find one that optimizes a chosen criterion function.

$$\{(X_i, y_i)\} \rightarrow \boxed{\begin{array}{c} \text{Learning Algorithm} \\ \text{(searching over } \mathcal{F}) \end{array}} \rightarrow f \in \mathcal{F}$$

- ▶ The question is: how can we formalize correctness of learning?
- ▶ There are many ways of addressing this issue (MDL, VC-theory etc).

- ▶ We discuss (at an elementary level) a specific statistical approach to address the issue of correctness of learning.
- ▶ We begin with a simple formalism where there is no 'noise' and the goal of learning is well-defined.

A Learning problem is defined by giving:

- (i) \mathcal{X} – input space; often \mathbb{R}^d (*feature space*)
- (ii) $\mathcal{Y} = \{0, 1\}$ – output space (*set of class labels*)
- (iii) $\mathcal{C} \subset 2^{\mathcal{X}}$ – concept space (*family of classifiers*)

Each $C \in \mathcal{C}$ is a subset of \mathcal{X} .

It can also be viewed as a function $C : \mathcal{X} \rightarrow \{0, 1\}$, with $C(X) = 1$ iff $X \in C$.

- (iv) $S = \{(X_i, y_i), i = 1, \dots, n\}$ – the set of examples,
 X_i are drawn *iid* according to some distribution P_x on \mathcal{X}
 $y_i = C^*(X_i)$ for some $C^* \in \mathcal{C}$.
 C^* is called target concept.

- ▶ We are considering a 2-class case.
- ▶ Hence any classifier is a function $C : \mathcal{X} \rightarrow \{0, 1\}$.
- ▶ Thus, \mathcal{C} is a family of classifiers.
- ▶ We call this concept space because we can say the system is learning a 'concept' from examples.
- ▶ The learning algorithm knows \mathcal{X} , \mathcal{Y} , \mathcal{C} ; but does not know C^* .
- ▶ It needs to learn the target concept from examples.

Some Comments

- ▶ We do not know the distribution P_x .
- ▶ We are trying to teach a concept through examples that come from an arbitrary distribution.
- ▶ However, taking that the examples are *iid* ensures we get 'representative' examples.
- ▶ Since we have taken $y_i = C^*(X_i)$, $\forall i$, there is no 'noise'.
- ▶ Also assuming that $C^* \in \mathcal{C}$ means that ideally we can learn the target concept. (Realizability assumption)

- ▶ We could take $\mathcal{C} = 2^{\mathcal{X}}$.
- ▶ This means we are searching over the family of all possible (2-class) classifiers.
- ▶ This may not be viable.
- ▶ We can choose a particular \mathcal{C} based on either some knowledge we have about the problem or because of the kind of learning algorithm we have.
- ▶ For example we can take \mathcal{C} to be all half-spaces – the family of all linear classifiers.

Probably Approximately Correct Learning

- ▶ Let us now try to define the goal of learning.
- ▶ Note that each $C \in \mathcal{C}$ can be viewed either as a subset of \mathcal{X} or a binary valued function on \mathcal{X} .
- ▶ Let C_n denote the concept or classifier output by the learning algorithm after it processes n *iid* examples.
- ▶ For correctness of the learning algorithm we want C_n to be 'close' to C^* as n becomes large.
- ▶ The closeness of C_n to C^* is in terms of classifying samples drawn from \mathcal{X} according to P_x .

- ▶ We define **error** of C_n by

$$\begin{aligned}\text{err}(C_n) &= P_x(C_n \Delta C^*) = P_x((\bar{C}_n \cap C^*) \cap (C_n \cap \bar{C}^*)) \\ &= \text{Prob}[\{X \in \mathcal{X} : C_n(X) \neq C^*(X)\}]\end{aligned}$$

- ▶ The $\text{err}(C_n)$ is the probability that on a random sample, drawn according to P_x , the classification of C_n and C^* differ.

- ▶ Essentially, we want $\text{err}(C_n)$ to become zero as $n \rightarrow \infty$.
- ▶ However, $\text{err}(C_n)$ is a random variable because C_n is a function of the random samples X_1, \dots, X_n .
- ▶ We take the above convergence to be in probability.

PAC learning

- ▶ We say a learning algorithm **Probably Approximately Correctly** (PAC) learns a concept class \mathcal{C} if given any $\epsilon, \delta > 0$, $\exists N < \infty$ such that

$$\text{Prob}[\text{err}(C_n) > \epsilon] < \delta$$

for all $n > N$ and for any distribution P_x and any C^* .

- ▶ The probability above is with respect to the distribution of n -tuples of *iid* samples drawn according to P_x on \mathcal{X} .
- ▶ The P_x is arbitrary. But, for testing and training the distribution is same – ‘fair’ to the algorithm.

- ▶ An algorithm PAC learns \mathcal{C} if

$$\text{Prob}[\text{err}(C_n) > \epsilon] < \delta$$


for sufficiently large n and any P_x .

- ▶ If $\text{err}(C_n) \leq \epsilon$, then C_n is ‘approximately correct’.
- ▶ So, what the above says is that the classifier output by the algorithm after seeing n random examples, C_n , is **approximately correct with a high probability**.
- ▶ The ϵ and δ are called the accuracy and confidence parameters respectively.

A Simple Example


- ▶ Suppose we want to learn the concept of medium-build persons based on features of height and weight.
- ▶ Here $\mathcal{X} = \mathbb{R}^2$.
- ▶ We would be given examples (with no errors!) drawn from some arbitrary distribution.

- The examples for learning our concept could be the following.



- ▶ What could be \mathcal{C} in this example?
- ▶ We can take \mathcal{C} to be all subsets of \mathbb{R}^2 .
- ▶ Or we can use some problem-based intuition and choose \mathcal{C} to be all axis-parallel rectangles.
- ▶ In this case, assuming $C^* \in \mathcal{C}$ means that the ‘god-given’ classifier is also an axis-parallel rectangle.

- The examples along with C^* now could be




- ▶ Now consider a very general class of algorithms.
- ▶ The strategy of the learning algorithm is as follows.
- ▶ The algorithm outputs a classifier which correctly classifies all examples.
- ▶ This is called a consistent classifier. Since $C^* \in \mathcal{C}$, we know there is always at least one consistent classifier.
- ▶ If there is more than one $C \in \mathcal{C}$ that is **consistent** with all examples, we output the 'smallest' such C .

- ▶ For finite sets, smallest is in terms of number of points; for other sets it is in terms of the 'areas' of the sets.
- ▶ We take \mathcal{C}_1 to be the set of all axis-parallel rectangles.
- ▶ We take \mathcal{C}_2 to be $2^{\mathcal{X}}$; that is, set of all possible classifiers.
- ▶ We will look at what we can say about the output of the algorithm in the two cases.

- ▶ We assume that C^* is an axis-parallel rectangle.
- ▶ Then C^* belongs to both \mathcal{C}_1 and \mathcal{C}_2 .
- ▶ All our examples are classified according to C^* .
- ▶ We assume that the boundary of the rectangle is part of C^* .

- ▶ First consider \mathcal{C}_1 .
- ▶ The smallest $C \in \mathcal{C}$ consistent with all examples would be the smallest axis-parallel rectangle enclosing all the positive examples seen so far.


- The concept output by the algorithm (when using \mathcal{C}_1) would be the following.



- ▶ Thus, under the strategy of our learning algorithm, for all n , the C_n would always be inside the C^* .
- ▶ Now let us show that this is a PAC learning algorithm.

- ▶ Whenever any example is classified as positive by C_n it would also be classified positive by C^* .
- ▶ Hence the points of \mathcal{X} where C_n makes errors is the annular region.

- ▶ Since C_n is also an axis-parallel rectangle which is inside C^* , the $C_n \Delta C^*$ would be the annular region between the two rectangles.



- ▶ Hence, $\text{err}(C_n)$ is the P_x -probability of this annular region.
- ▶ Note that we are not really bothered about the area of this annular region; we are only interested in the probability mass of this region under P_x .

- ▶ Now, given an $\epsilon > 0$, we have to bound the probability that $\text{err}(C_n) > \epsilon$.
- ▶ The error is greater than ϵ only if the probability mass (under P_x) of the annular region is greater than ϵ .
- ▶ When does this event $[\text{err}(C_n) > \epsilon]$ occur?
- ▶ Only when none of the examples seen happen to be in the annular region.
- ▶ Why? – Otherwise, the rectangle learnt by our algorithm would have been closer to C^* .

- ▶ Hence the probability of the event $[\text{err}(C_n) > \epsilon]$ is same as the probability that when n *iid* examples are drawn according to P_x none of them came from a subset of \mathcal{X} that has P_x -probability at least ϵ .
- ▶ That is, all examples came from a subset of probability at most $(1 - \epsilon)$.
- ▶ The probability of this happening is at most $(1 - \epsilon)^n$.

- ▶ Hence we have

$$\text{Prob}[\text{err}(C_n) > \epsilon] \leq (1 - \epsilon)^n$$

- ▶ Let N be such that $(1 - \epsilon)^n < \delta$, for all $n > N$.
- ▶ The required N is $N \geq \frac{\ln(\delta)}{\ln(1-\epsilon)}$
(bound on number of examples).
- ▶ For this N , we have

$$\text{Prob}[\text{err}(C_n) > \epsilon] \leq \delta, \forall n \geq N$$

showing the algorithm PAC learns the concept class.

- ▶ Now let us consider the same algorithm with concept class $\mathcal{C}_2 = 2^{\mathcal{X}}$.
- ▶ Here we are searching over all possible 2-class classifiers.
- ▶ So, intuitively, we do not expect the algorithm to be able to learn anything.
- ▶ There is too much ‘flexibility’ in the bag of classifiers over which we are searching.
- ▶ Let us show this formally.

- ▶ What would be C_n now?
- ▶ After seeing n examples, the smallest set in \mathcal{C}_2 that is consistent with all examples is the set consisting of all the positive examples seen so far!!
- ▶ Now the algorithm simply remembers all the positive examples seen.
- ▶ This happened because every possible finite subset of \mathcal{X} is in our concept class.

- ▶ So, now, $C_n \Delta C^*$ would be the axis parallel rectangle C^* minus some finite number of points from it.
- ▶ So, under any continuous P_x ,
 $\text{err}(C_n) = P_x(C_n \Delta C^*) = P_x(C^*)$.
- ▶ Hence, for $\epsilon < P_x(C^*)$, $\text{Prob}[\text{err}(C_n) > \epsilon] = 1$ for all n .
- ▶ Thus, the algorithm can not PAC learn with \mathcal{C}_2 .

- ▶ This example clearly illustrates the difficulty of learning from examples if the bag of classifiers being considered is too 'large'.
- ▶ The largeness is not in terms of number of elements in our concept class.
- ▶ Both \mathcal{C}_1 and \mathcal{C}_2 contain uncountably infinite number of classifiers.
- ▶ We would later define an appropriate quantity to quantify the sense in which one concept class can be said to be bigger than (or more complex than) another.

- ▶ At this point we can still see how \mathcal{C}_1 is smaller than \mathcal{C}_2 .
- ▶ Since every axis parallel rectangle can be specified by four quantities, this class can be parameterized by four parameters.
- ▶ However, there is no such finite parameterization for $\mathcal{C}_2 = 2^{\mathbb{R}^2}$.
- ▶ Also, the strategy of our algorithm can be coded efficiently in case of \mathcal{C}_1 .

- ▶ The concept of PAC learnability is interesting.
- ▶ It allows one to properly define what is correctness of learning and allows us to ask questions like whether a given algorithm learns correctly.
- ▶ As we have seen in our example, we can also bound the number of examples needed to learn to a given level of accuracy and confidence.
- ▶ Thus we can appreciate relative complexities of different learning problems.

- ▶ However, PAC learnability deals with ideal learning situations.
- ▶ We assume there is a ('god-given') C^* and that it is in our \mathcal{C} .
- ▶ Also, we assume that examples are noise free and are perfectly classified.
- ▶ Next we consider an extension of this framework that is relevant for realistic learning scenarios.

In our new framework we are given

- ▶ \mathcal{X} – input space; (as earlier, *Feature space*)
- ▶ \mathcal{Y} – Output space (as earlier, *Set of class labels*)
- ▶ \mathcal{H} – hypothesis space (*family of classifiers*)

Each $h \in \mathcal{H}$ is a function: $h : \mathcal{X} \rightarrow \mathcal{A}$

where \mathcal{A} is called *action space*.

- ▶ Training data: $\{(X_i, y_i), i = 1, \dots, n\}$
drawn *iid* according to some distribution P_{xy} on $\mathcal{X} \times \mathcal{Y}$.

Some Comments

- ▶ We have replaced \mathcal{C} with \mathcal{H} .
- ▶ If we take $\mathcal{A} = \mathcal{Y}$ then it is same as earlier.
- ▶ But the freedom in choosing \mathcal{A} allows for taking care of many situations.
- ▶ For example, when $\mathcal{Y} = \{0, 1\}$, we can take $\mathcal{A} = [0, 1]$ (e.g., logistic regression).

- ▶ Now we draw examples from $\mathcal{X} \times \mathcal{Y}$ according to P_{xy} . This allows for 'noise' in the training data.
- ▶ For example, when class conditional densities overlap, same X can come from different classes with different probabilities.
- ▶ We can always factorize $P_{xy} = P_x P_{y|x}$. If $P_{y|x}$ is a degenerate distribution then it will be same as earlier – we draw iid samples from \mathcal{X} and each point is essentially classified by the target classifier.
- ▶ However, having examples drawn from $\mathcal{X} \times \mathcal{Y}$ using a distribution, allows for many more scenarios.

- ▶ As before, the learning machine outputs a hypothesis, $h_n \in \mathcal{H}$, given the training data consisting of n examples.
- ▶ However, now there is no notion of a target concept/hypothesis.
- ▶ There may be no $h \in \mathcal{H}$ which is consistent with all examples.
- ▶ Hence we use the idea of loss functions to define the goal of learning.

Loss function

- ▶ Loss function: $L : \mathcal{Y} \times \mathcal{A} \rightarrow \mathbb{R}^+$.
- ▶ The idea is that $L(y, h(X))$ is the 'loss' suffered by $h \in \mathcal{H}$ on a (random) sample $(X, y) \in \mathcal{X} \times \mathcal{Y}$.
- ▶ More generally we can let loss depend on X also explicitly and can write $L(X, y, h(X))$ for loss function.
- ▶ By convention we assume that the loss function is non-negative.
- ▶ Now we can look for hypotheses that have low average loss over samples drawn according to P_{xy} .

Risk Function

- ▶ Define the **risk** function, $R : \mathcal{H} \rightarrow \mathbb{R}^+$, by

$$R(h) = E[L(y, h(X))] = \int L(y, h(X)) dP_{xy}$$

- ▶ Risk is expectation of loss where expectation is with respect to P_{xy} .
- ▶ We want to find h with low risk.

Risk Minimization

- ▶ Let

$$h^* = \arg \min_{h \in \mathcal{H}} R(h)$$

- ▶ We define the goal of learning as finding h^* , the global minimizer of risk.
- ▶ Risk minimization is a very general strategy adopted by most machine learning algorithms.
- ▶ Sometimes called 'Agnostic Learning'
- ▶ Note that we may not have any knowledge of P_{xy} .
- ▶ Minimization of $R(\cdot)$ directly is not feasible.

Empirical Risk function

- ▶ Define the **empirical risk function**, $\hat{R}_n : \mathcal{H} \rightarrow \mathbb{R}^+$, by

$$\hat{R}_n(h) = \frac{1}{n} \sum_{i=1}^n L(y_i, h(X_i))$$

This is the sample mean estimator of risk obtained from n *iid* samples.

- ▶ Let \hat{h}_n^* be the global minimizer of empirical risk, \hat{R}_n .

$$\hat{h}_n^* = \arg \min_{h \in \mathcal{H}} \hat{R}_n(h)$$

Empirical Risk Minimization

- ▶ Given any h we can calculate $\hat{R}_n(h)$.
- ▶ Hence, we can (in principle) find \hat{h}_n^* by optimization methods.
- ▶ Approximating h^* by \hat{h}_n^* is the basic idea of empirical risk minimization strategy.
- ▶ Used in most ML algorithms.

- ▶ Is \hat{h}_n^* a good approximator of h^* , the minimizer of true risk (for large n)?
- ▶ This is the question of **consistency of empirical risk minimization**.
- ▶ Thus, we can say a learning problem has two parts.
 - ▶ The optimization part: find \hat{h}_n^* , the minimizer of \hat{R}_n .
 - ▶ The statistical part: Is \hat{h}_n^* a good approximator of h^* .

- ▶ The optimization part depends on the loss function.
- ▶ Note that the loss function is chosen by us; it is part of the specification of the learning problem.
- ▶ The loss function is intended to capture how we would like to evaluate performance of the classifier.
- ▶ We look at a few loss functions in the 2-class case.

The 0–1 loss function

- ▶ Let $\mathcal{Y} = \{0, 1\}$ and $\mathcal{A} = \mathcal{Y}$.
- ▶ Now, the 0–1 loss function is defined by

$$L(y, h(X)) = I_{[y \neq h(X)]}$$

where $I_{[A]}$ denotes indicator of event A .

- ▶ The 0-1 loss function is

$$L(y, h(X)) = I_{[y \neq h(X)]}$$

- ▶ Risk is expectation of loss.
- ▶ Hence, $R(h) = \text{Prob}[y \neq h(X)]$;
the risk is probability of misclassification.
- ▶ So, h^* minimizes probability of misclassification.
(Bayes classifier)

- ▶ Here we assumed that the learning algorithm searches over a class of binary-valued functions on \mathcal{X} .
- ▶ We can extend this to, e.g., discriminant function learning.
- ▶ We take $\mathcal{A} = \mathfrak{R}$ (now $h(X)$ is a discriminant function).
- ▶ We can define the 0-1 loss now as

$$L(y, h(X)) = I_{[y \neq \text{sgn}(h(X))]}$$

- ▶ Having any fixed misclassification costs is essentially same as 0–1 loss.
- ▶ Even if we take $\mathcal{A} = \Re$, the 0–1 loss compares only sign of $h(x)$ with y . The magnitude of $h(x)$ has no effect on the loss.
- ▶ Here, we can not trade ‘good’ performance on some data with ‘bad’ performance on others.
- ▶ This makes 0–1 loss function more robust to noise in classification labels.

- ▶ While 0–1 loss is an intuitively appealing performance measure, minimizing empirical risk here is hard.
- ▶ The 0–1 loss function is non-differentiable which makes the empirical risk function also non-differentiable.
- ▶ Hence many other loss functions are often used in Machine Learning.

Squared error loss

- ▶ The squared error loss function is defined by

$$L(y, h(X)) = (y - h(X))^2$$

- ▶ As is easy to see, the linear least squares method is empirical risk minimization with squared error loss function.
- ▶ Here we can take \mathcal{Y} as $\{0, 1\}$ or $\{+1, -1\}$. We take $\mathcal{A} = \mathfrak{R}$ so that each h is a discriminant function.
- ▶ As we know, we can use this for regression problems also and then we take $\mathcal{Y} = \mathfrak{R}$.

- ▶ Another interesting scenario here is to take $\mathcal{Y} = \{0, 1\}$ and $\mathcal{A} = [0, 1]$.
- ▶ Then each h can be interpreted as a posterior probability (of class-1) function.
- ▶ As we know, the minimizer of expectation of squared error loss (the risk here) is the posterior probability function.
- ▶ So, risk minimization would now look for a function in \mathcal{H} that is a good approximation for the posterior probability function.

- ▶ The empirical risk minimization under squared error loss is a convex optimization problem for linear models (when h is linear in its parameters).
- ▶ The squared error loss is extensively used in many learning algorithms.

soft margin loss or hinge loss

- ▶ Take $\mathcal{Y} = \{+1, -1\}$ and $\mathcal{A} = \mathbb{R}$. The loss function is given by

$$L(y, h(X)) = \max(0, 1 - yh(X))$$

- ▶ Here, if $yh(X) > 0$ then classification is correct and if $yh(X) \geq 1$, loss is zero.
- ▶ This also results in convex optimization for empirical risk minimization.

Margin Losses

- ▶ All three losses we mentioned can be written as function of $yh(X)$ by taking $\mathcal{Y} = \{-1, +1\}$.
- ▶ The 0-1 loss :

$$L(y, h(X)) = \text{sign}(-yh(X))$$


- ▶ The squared error loss:

$$L(y, h(X)) = (y - h(X))^2 = (1 - yh(X))^2$$

- ▶ The hinge loss (used in SVM):

$$L(y, h(X)) = \max(0, 1 - yh(X))$$

Plot of 2-class loss functions



- We can think of the other losses as convex approximations of 0-1 loss.

- ▶ As we saw, there are many different loss functions one can think of.
- ▶ Many of them also make the empirical risk minimization problem efficiently solvable.
- ▶ We consider many such algorithms in this course.
- ▶ Now, let us get back to the statistical question that we started with.

Consistency of Empirical Risk Minimization

- ▶ Our objective is to find h^* , minimizer of risk $R(\cdot)$.
- ▶ We minimize the empirical risk, \hat{R}_n , and thus find \hat{h}_n^* .
- ▶ We want h^* and \hat{h}_n^* to be 'close'.
- ▶ More precisely we are interested in the question: Does

$$\forall \delta > 0, \text{ Prob}[|R(\hat{h}_n^*) - R(h^*)| > \delta] \rightarrow 0, \text{ as } n \rightarrow \infty?$$

- ▶ Same as asking whether $R(\hat{h}_n^*)$ converges in probability to $R(h^*)$
- ▶ This is the question we will address in the next class.