

# Recap – Some Issues in learning neural networks

- ▶ Activation function for nodes
- ▶ Online or batch mode for learning: fixing minibatch size
- ▶ Normalization of inputs
  - ▶ convert individual features to mean zero variance 1,
  - ▶ whitening transform
- ▶ Learning algorithm
  - ▶ BP with momentum
  - ▶ ADAM algorithm
  - ▶ Weight decay, dropout, step-size selection
  - ▶ other optimization methods (e.g., Hessian based)
- ▶ Loss functions (e.g., cross entropy loss)
- ▶ Network structure
  - ▶ Number of hidden layers and hidden nodes
  - ▶ The initial values for weights

# Recap


- ▶ Over the last 5-10 years there is increasing interest in deep networks – networks with large number of hidden layers.
- ▶ Deep networks are shown to be highly effective in applications.
- ▶ Learning deep networks through gradient descent can be challenging because of the large number of weights.
- ▶ One approach is to use 'local' connectivity patterns (e.g., CNNs)
- ▶ Another approach is to have good data-dependent initializations of weights.

# Unsupervised Learning to Initialize Weights

- ▶ One of the insights that contributed to the successes of learning of deep networks is that we need good initialization **based on data**.
- ▶ One possible way to realize deep learning is to obtain good initial weights through unsupervised learning.
- ▶ Two important developments here gave thrust to deep learning:
  - ▶ Autoencoders to initialize multilayer feedforward nets.
  - ▶ Restricted Boltzman machines to initialize Deep belief networks.


# Autoencoder Network

- ▶ It is a feedforward network that was investigated in 1980's.
- ▶ It is a 3-layer feedforward network with very few hidden nodes.




$$Z = f(WX + b)$$
$$X' = f(W^T X + b')$$

# Autoencoder Network



- ▶ We want to learn  $W$  such that  $X' = X$ .
- ▶ The  $W$  is good for transforming the representation.
- ▶ The motivation is from Minimum Description Length principle

# Autoencoder Network



- ▶ Given data,  $\{X^1, \dots, X^m\}$ , we can learn  $W$  through backpropagation.
- ▶ This is unsupervised learning.
- ▶ We learn a 'compressed' representation.
- ▶ The representation may be 'good' because we can recreate original  $X$ .
- ▶ Like dimensionality reduction. (Using  $W$  and  $W^T$  is motivated by the linear case)

- ▶ There are many variations possible on this basic theme.

# Denoising Autoencoder

- ▶ We would give noise-corrupted  $X$  at input but want  $X'$  to be the 'clean'  $X$ .
- ▶ Can add independent noise to each component of  $X$ .
- ▶ But, what is often done is to make a few randomly selected components of  $X$  zero.
- ▶ If we can learn  $W$  to create  $X$  at output, then that  $W$  can capture dependences among components of  $X$ .
- ▶ Hence,  $W$  is a good set of weights to transform  $X$  into a useful representation.



## Another Variation: Sparse Autoencoder

- ▶ We are learning a 'compressed' representation by having only few nodes in hidden layer.
- ▶ Alternately, only a few of the hidden nodes should be 'active' for any given  $X$ .
- ▶ Then we need not have any restriction on the number of hidden layer nodes
- ▶ By making representation 'sparse' we achieve similar 'compression'
- ▶ Sparsity can be incorporated into the objective function.

# Ensuring Sparsity

- ▶ Let  $X^i$  be the  $i^{th}$  example,  $i = 1, \dots, m$ . Let

$$\hat{\rho}_j = \frac{1}{m} \sum_{i=1}^m y_j^2(X^i)$$

( $y_j^2$  is the output of  $j^{th}$  node in the hidden layer)

- ▶  $\hat{\rho}_j$  is the fraction of input for which  $j^{th}$  hidden neuron is 'ON'.
- ▶ We want  $\hat{\rho}_j = \rho$ ,  $\forall j$ , where  $\rho$  is the sparsity parameter.
- ▶ Typically  $\rho$  is very small (e.g., 0.05).


- ▶ We can think of  $\rho$  as a bernoulli parameter.
- ▶ We can use KL divergence to measure deviation of  $\rho_j$  from  $\rho$ .

$$KL(\rho||\rho_j) = \rho \log \left( \frac{\rho}{\rho_j} \right) + (1 - \rho) \log \left( \frac{1 - \rho}{1 - \rho_j} \right)$$

- ▶ We can use the following objective function

$$J(W) = \sum_{i=1}^m \|y^3(X^i) - X^i\|^2 + \beta \sum_{j=1}^{n_2} KL(\rho||\rho_j)$$

$$J = \sum_{i=1}^m \|y^3(X^i) - X^i\|^2 + \beta \sum_{j=1}^{n_2} \rho \log \left( \frac{\rho}{\rho_j} \right) + (1 - \rho) \log \left( \frac{1 - \rho}{1 - \rho_j} \right)$$



- ▶ The  $\rho_j$  depend on  $y_j^2$  but not on  $w_{ij}^2$ .
- ▶ Hence  $\delta_j^3$  would be same as earlier.
- ▶ We only need to calculate  $\frac{\partial J}{\partial w_{ij}^1}$ .

- It can be shown that


$$\delta_i^2 = \left( \sum_{j=1}^{n_3} w_{ij}^2 \delta_j^3 + \beta \left( \frac{-\rho}{\rho_i} + \frac{1-\rho}{1-\rho_i} \right) \right) f'(\eta_i^2)$$

- Note that we have  $\frac{\partial J}{\partial w_{ij}^1} = \delta_j^2 y_i^1$ .
- We need current value of  $\rho_j$  for weight update – need  $y_j^2(X^i)$  for all  $i$ .
- Need one pass over examples before starting weight update.


- ▶ We can have the sparsity penalty in a denoising autoencoder too.
- ▶ So, we get sparse denoising autoencoders.
- ▶ We can use the autoencoder network to initialize weights successively in each layer of a feedforward network.

# Autoencoder for initializing weights

- ▶ We will first learn the weights from input layer to first hidden layer using (sparse denoising) autoencoder. (This is unsupervised learning)




We learn the first layer weights as below:



- ▶ Here we have no control on the number of hidden nodes in the autoencoder network,
- ▶ Hence we use sparse autoencoder.




Next we learn the weights from first hidden layer to second hidden layer using autoencoder.



- Input to autoencoder are training data  $X^i$  transformed through the  $W$  that is already learnt.

Following this procedure we learn weights for each successive layer.



- ▶ The basic idea here is the following.
- ▶ Using one hidden layer at a time we learn weights to generate a new representation of the input.
- ▶ The weights are learnt to recreate the input through a sparse autoencoder.
- ▶ Such weights form a good (and data-dependent) initialization of weights in the network.
- ▶ Starting with this initialization, we use supervised learning through backpropagation to learn all the weights in the network.
- ▶ This idea of using autoencoder for initializing all weights for deep network is very effective.
- ▶ In many applications this allows backpropagation to learn 'good local minima' even for very deep networks.

# Convolutional Neural Networks (CNNs)


- ▶ CNNs represent a deep network model that has revolutionized image-recognition
- ▶ CNNs are largely responsible for starting the current wave of interest in deep neural networks.

# Convolutional Neural Networks (CNNs)

- ▶ CNNs are deep neural networks that are originally proposed for image recognition.
- ▶ The input layer of a CNN is 2-dimensional because it is an image.
- ▶ There are many features of CNNs that make them much more efficient compared to normal feedforward nets for image based pattern recognition.
- ▶ They use local connectivity, weight sharing, multiple 'feature planes' to learn appropriate features from data.
- ▶ We start by looking at the ideas underlying CNNs first in a one dimensional context.

# Connectivity is local


- ▶ In fully connected networks weights in a layer grow as square of number of nodes.
- ▶ We can reduce the number of weights by making connections local.



- ▶ Now weights scale only linearly with number of neurons.


# Weight-Sharing

- ▶ We can reduce weights further by 'sharing' of weights.



- ▶ Now number of weights per layer is a constant independent of the number of nodes.
- ▶ Such layers are called convolutional layers.

# Output of a convolutional node



- ▶ The output can be calculated as

$$y_2 = f(w_{-1}x_1 + w_0x_2 + w_{+1}x_3)$$


$$y_3 = f(w_{-1}x_2 + w_0x_3 + w_{+1}x_4)$$

- ▶ Essentially a convolution type computation



# Output of a convolutional node

- ▶ The general case.



$$y_i = f \left( \sum_{k=-1}^{+1} w_k x_{i+k} \right)$$

- ▶ Weight vector is a 3-point filter.
- ▶ We are computing the filter output at each point by sliding over the input.

# CNNs

- ▶ Each layer of a network with such connectivity is called a convolutional layer.
- ▶ Convolutional neural networks (CNNs) are feedforward networks that contain many such convolutional layers.

- ▶ In a CNN, the output of the convolutional layer is passed through a non-linear activation function.
- ▶ The often used activation function is ReLU.

# CNNs

- ▶ What are typical problems where such connectivity is natural?
- ▶ Example: Image-based pattern recognition
- ▶ Useful features in an image (e.g., edges, corners) are computed using such local convolution through so called masks.
- ▶ We apply the same operation at all points in an image to detect the feature wherever it exists.

- ▶ For example, this is a simple edge-detector mask

-1	0	1
-1	0	1
-1	0	1

- ▶ We do this masking (or convolution) operation at each point in the image

- ▶ Traditionally, in Pattern Recognition, feature extraction and classification are viewed as two separate steps.
- ▶ Often features are designed separately based on the knowledge of the problem.
- ▶ For example, the SIFT or HOG features used in image recognition problems.
- ▶ After transforming the input image into a representation using the chosen features, one learns a (linear or nonlinear) classifier.
- ▶ As discussed earlier, the philosophy underlying the neural networks approach is that we should **automatically learn** the relevant features based on the data.

- ▶ Each convolutional layer is essentially detecting a feature.
- ▶ Since the weights would be learnt, we are learning the 'proper features' automatically using the training data.
- ▶ So far we are considering 1D layers.
- ▶ Image is two dimensional and hence all layers as well as the filters need to be two dimensional.
- ▶ But this is a straight-forward extension as we shall see.


# Multiple Convolutional Layers

- ▶ We need to combine simple features into more complex features to achieve object recognition.
- ▶ We may combine edge pixels at different locations into lines before we can recognize shapes.
- ▶ This is easily achieved by having many convolutional layers.




# Receptive Field of a convolutional layer node

- ▶ The second convolutional layer gets its output from the first, through local connectivity.
- ▶ Eventhough connectivity is local, later filters are effectively looking at larger portion of the input.  
(We do not explicitly show the nonlinear activation)



# Receptive Field of a convolutional layer node


- ▶ The second convolutional layer gets its output from the first, through local connectivity.
- ▶ Eventhough connectivity is local, later filters are effectively looking at larger portion of the input.  
(We do not explicitly show the nonlinear activation)



- ▶ Each node in any convolutional layer has an effective receptive field in the input.

# Output of a Convolutional Layer

- ▶ Since we have many layers let us revert to earlier notation.



$$y_i^l = f \left( \sum_{k=-1}^{+1} w_k^l y_{i+k}^{l-1} \right)$$

where  $w_k^l$  is the weight connecting any node in layer  $l$  to a node with offset  $k$  in layer  $l - 1$ .

- ▶ When we generalize to 2-D case,  $y_i^l$  would become  $y_{ij}^l$ ,  $w_k^l$  would become  $w_{sk}^l$  and so on.


- ▶ The convolutional layer we saw so far is completely specified by one weight vector (of small dimension).
- ▶ It represents a single feature detector or filter.
- ▶ Our convolutional layer essentially detects this feature wherever it exists in the image.
- ▶ The output of the convolutional layer can be thought of as the representation of the input in terms of this feature (feature plane).

# Multiple Filters

- ▶ A single feature would not be sufficient for many pattern recognition tasks.
- ▶ We would need multiple filters.
- ▶ For example, at each point we may want to detect edges in different orientations.
- ▶ Further, through multiple layers we need to combine multiple simple features into multiple complex features.
- ▶ Hence, every convolutional layer should have multiple filters.


# Convolutional layer with multiple filters

- ▶ As earlier let us take input as 1D.
- ▶ Now the convolutional layer would be 2D because of multiple filters



# Multiple convolutional layers with multiple filters

- ▶ Consider two convolutional layers with multiple filters.
- ▶ Each layer is 2D (space dimension and filter dimension)



- ▶ Connectivity is full in the filter dimension

# Notation for multiple filters

- ▶ Each output would now have three 'indices' – layer, position in the layer, and filter.
- ▶  $y_i^{l,m}$  – output of node- $i$ , layer- $l$ , filter- $m$ .
- ▶ The connectivity is local in space but is full in filter domain (it needs to combine all features at that point).
- ▶ Each weight has four 'indices' – layer and filter number in that layer; the space-offset and filter coordinates of the input it is multiplying.
- ▶  $w_{k,m'}^{l,m}$  – weight for filter  $m$  in layer  $l$  connecting to node with 'offset'  $k$  in space and filter  $m'$  in layer  $l - 1$ .



# Output of a convolutional layer node with multiple features

- ▶  $y_i^{l,m}$  – output of node- $i$ , layer- $l$ , filter- $m$ .
- ▶  $w_{k,m'}^{l,m}$  – weight for filter  $m$  in layer  $l$  connecting to node with ‘offset’  $k$  in space and filter  $m'$  in layer  $l - 1$ .
- ▶ The outputs are now calculated as

$$y_i^{l,m} = f \left( \sum_{m'} \sum_{k=-q}^q w_{k,m'}^{l,m} y_{i+k}^{l-1,m'} \right)$$

$W^{l,m}$  – the ‘weight matrix’ associated with filter  $m$  in layer  $l$ .


- ▶ Compare with the single filter case

$$y_i^l = f \left( \sum_{k=-1}^{+1} w_k^l y_{i+k}^{l-1} \right)$$

# Notation for multiple filters

- ▶ Each output now has three 'indices' – layer, position in the layer, and filter.
- ▶ Each weight has four 'indices' – layer and filter number in that layer; the space-offset and filter coordinates of the input it is multiplying.
- ▶ Essentially now all these are hyper-matrices (also called tensors).
- ▶ When we move to 2D, the space part would be a pair of coordinates / offsets.


- ▶ Consider a single filter dimension.
- ▶ As we said earlier, there is local connectivity and weight sharing. That is what defines the filter



- ▶ There is one more characteristic for the filter definition.



# Stride of a filter

- ▶ In the connectivity there can be an 'off-set'
- ▶ The filter shown below is said to have stride 2.




# Stride of a filter

Filter with stride 1 and stride 2.



# Zero Padding

- ▶ In all the figures shown so far, successive layers seem to have fewer nodes.
- ▶ We need not have that. Can use 'zero-padding'.



# Numbers of nodes in successive layers



- ▶ In general, we can have different numbers in successive layers.
- ▶ Suppose we have  $n$  nodes in one layer and use  $P$  number of nodes for zero padding on either side. Suppose filter width is  $F$  and its **stride** is  $S$ .
- ▶ Then every node but the last one in the next layer would have  $S$  nodes exclusively to it and last node would have  $F$ . If this layer has  $n'$  nodes, then

$$(n' - 1)S + F = n + 2P$$

# Numbers of nodes in successive layers

- Suppose  $n$  and  $n'$  are nodes in successive layers,  $P$  is number of nodes for zero padding on either side, filter width is  $F$  and its **stride** is  $S$ .

$$(n' - 1)S + F = n + 2P$$





# Numbers of nodes in successive layers


- ▶ Suppose we have  $n$  nodes in one layer and use  $P$  number of nodes for zero padding on either side. Suppose filter width is  $F$  and its **stride** is  $S$ . Then, number of nodes in the next layer,  $n'$  satisfies

$$(n' - 1)S + F = n + 2P$$

- ▶ Note that all quantities here have to be integers.
- ▶ If you do not use zero padding, number of nodes in successive layers decreases.
- ▶ Even otherwise, we may want sizes of convolutional layers successively reduced.

# Pooling layers


- ▶ One also reduces the size of the successive convolutional layers by using what is known as pooling.
- ▶ For example, we can reduce size to half by taking average or Max of successive elements.



- ▶ Often, Max pooling is used.


- ▶ Each convolutional layer has many filters and each extracts some feature.
- ▶ Through a series of convolutional layers, the original input (image) is transformed to a new feature representation.
- ▶ Then we need a classifier to classify it.
- ▶ For this, we have one or more 'fully-connected' layers after all convolutional layers.
- ▶ Finally, we may have a soft-max layer for multi-class classification.

# A Typical CNN



- ▶ We need to extend all our notation to the case of images which are two dimensional in space.
- ▶ Each convolutional layer can be thought of as three dimensional – two space dimensions and one filter dimension.
- ▶ So, we can think of it as processing a ‘volume’ and producing a ‘volume’
- ▶ To keep notation uniform, we can think of input image also as 3-dimensional – e.g., colour being the third dimension.
- ▶ The third dimension in each layer, the **filters**, are also called the **channels**.

- ▶ Let us slightly modify our notation for better readability.
- ▶  $y_r^\ell(i, j)$  – the output of the node at spatial location  $(i, j)$  corresponding to filter  $r$  in (convolutional) layer  $\ell$ .
- ▶ Now the node corresponding to filter  $r$  in layer  $\ell$  would be associated with a weight tensor that connects this node to a limited spatial region and all filters in the previous layer. Let us look at 1D case:



- ▶ Let us slightly modify our notation for better readability.
- ▶  $y_r^\ell(i, j)$  – the output of the node at spatial location  $(i, j)$  corresponding to filter  $r$  in (convolutional) layer  $\ell$ .
- ▶ Now the node corresponding to filter  $r$  in layer  $\ell$  would be associated with a weight tensor that connects this node to a limited spatial region and all filters in the previous layer.
- ▶  $W_r^\ell(a, b; c)$  – weight connecting any node in layer  $\ell$  and filter  $r$  to a node in the previous layer at a spatial offset given by  $(a, b)$  and filter index in the previous layer,  $c$ .

- ▶ Now we can write the expression for output of a convolutional layer:

$$y_r^\ell(i, j) = \sum_c \sum_{a=-q}^q \sum_{b=-q}^q y_c^{\ell-1}(i+a, j+b) W_r^\ell(a, b; c)$$

where  $c$  ranges over filters in layer  $\ell - 1$ .

- ▶ In the above,  $r$  ranges over number of filters in that layer and  $i, j$  range over the spatial extent of that layer.
- ▶ Here we have taken the spatial offset to go from  $-q$  to  $q$ . (A notation)
- ▶ Hence filter size is  $(2q + 1) \times (2q + 1)$ .



- ▶ As we discussed earlier, after the convolution operation, we actually pass the output through a non-linearity.
- ▶ Suppose after the nonlinear activation function the output is  $\bar{y}_r^\ell(i, j)$ .
- ▶ Now the equations become

$$y_r^\ell(i, j) = \sum_c \sum_{a=-q}^q \sum_{b=-q}^q \bar{y}_c^{\ell-1}(i+a, j+b) W_r^\ell(a, b; c)$$

$$\bar{y}_r^\ell(i, j) = f(y_r^\ell(i, j)) = \max\{0, y_r^\ell(i, j)\}$$

where we are assuming ReLU activation.

- ▶ There are many important details that we have ignored in writing these equations.
- ▶ Recall that the number of nodes in successive layers are related by

$$(n_2 - 1) * S + F = n_1 + 2P$$

(This holds for both space dimensions)

- ▶ In our equations we assumed  $S = 1$ .
- ▶ Our convention of taking  $F = 2q + 1$  and letting the offset variable range over  $-q$  to  $q$  is convenient only when we have  $n_2 = n_1$  by taking  $P$  as large as needed.
- ▶ In such cases, we simply take the value to be zero when index is negative.

- ▶ Numbers of nodes are related by

$$(n_2 - 1) * S + F = n_1 + 2P$$

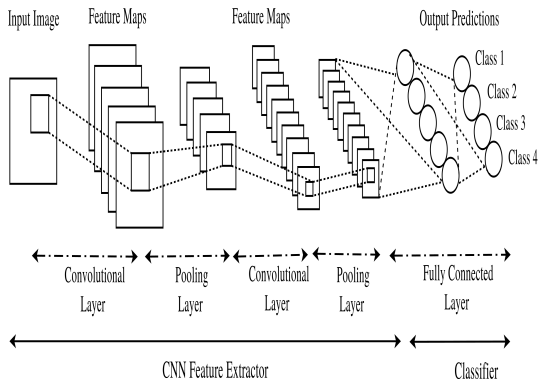
- ▶ Another special case is  $S = 1$  and  $P = 0$ . Then  $n_1 = n_2 - 1 + F$
- ▶ This would reduce the spatial size of the next convolutional layer
- ▶ Then we number nodes as 1 to  $n_1$  and 1 to  $n_2$ ; and let the offset variable run from 0 to  $F - 1$ .

- ▶ Our convention for numbering nodes as well as the 'offset' variables depends on how the number of nodes in successive layers changes.
- ▶ It also depends on the stride of the filter.
- ▶ We also need to properly take care of pooling layers.
- ▶ In general, writing these equations is a little more involved compared to those for regular feedforward networks.

- ▶ After a number of such convolutional layers, we have fully connected layers.
- ▶ This part is like a standard feedforward network we considered earlier.
- ▶ The input to the first fully connected layer would be a 'vectorized version' of the last convolutional layer.
- ▶ If the output of a convolutional layer needs to be an image, then we do not have the fully connected layers. ("Fully Convolutional Networks")

# A Typical CNN Architecture

- Thus, we have the following structure for a CNN.



- ▶ Fixing a CNN architecture involves many issues.
- ▶ Each convolutional layer is characterized by number of filters, size of filters and stride.
- ▶ Each such layer has one weight tensor.
- ▶ The spacial extents of different convolutional layers is now determined by whether or not to we use zero-padding.
- ▶ We also need to fix details of fully connected layers.
- ▶ There are many standard architectures (e.g., Alexnet, VGGNet, Resnet etc).

# Training a CNN


- ▶ To use a CNN as a classifier, we need to learn all weights.
- ▶ One normally uses the cross entropy loss.
- ▶ We need to learn all the filters (weights in convolutional layers) in addition to weights in fully connected layers.
- ▶ Fully connected layers are same as the earlier feedforward networks.



# Training a CNN


- ▶ We can learn the weights in convolutional layers (i.e., the filters) also using the backpropagation algorithm.
- ▶ Since convolutional layers do only linear summation followed by a nonlinear transformation, the backpropagation is essentially same as what we derived earlier.
- ▶ There is just one simple modification.

- ▶ In a CNN weights are shared.



- ▶ The final update for each weight is sum of all the updates.

- We can similarly backpropagate through pooling layer also.



- In case of max-pooling, we transfer the 'error at the node' to that node which was the maximum.

# Regularization for learning CNNs

- ▶ Many different regularizations are used.
- ▶ We have discussed weight decay and  $L_2$  regularization.
- ▶ Often also implemented as a constraint on norm of each weight vector
- ▶ Essentially, large values for weights is an indicator of overfitting.

# dropout

- ▶ We had mentioned dropout earlier
- ▶ In dropout regularization, one keeps dropping a random subset of nodes in the network from being considered.
- ▶ In each iteration, for each node (except possibly for the input nodes) we independently decide with probability,  $p$ , whether that node (along with all its incoming and outgoing links) would be present.
- ▶ The backpropagation would update only those links that are present.
- ▶ Another variant is dropconnect.

- ▶ For a general network, dropout can be represented as:

$$y_j^l = f(\eta_j^l) = f\left(\sum_i w_{ij}^{l-1} \xi_i^{l-1} y_i^{l-1}\right)$$

where  $\xi_i^{l-1} \sim \text{Bernouli}(p)$ .

- ▶ This can be particularized for any network structure, e.g., CNN
- ▶ In the equations, output of a node is multiplied by a Bernoulli random variable.
- ▶ The  $\xi_j^l$  are independent.
- ▶ This is effective in guarding against spurious 'co-adaptation' of weights.
- ▶ It essentially 'averages' many low-complexity networks and hence is effective as a regularization technique.

- ▶ Consider a single logistic unit (in some layer) with  $n$  inputs. Because of dropout it gets different subsets of inputs with different probabilities.
- ▶ Let  $S_1, \dots, S_m$  be the net input to this unit under these different subnetworks with probabilities  $P_1, \dots, P_m$  and let  $O_1, \dots, O_m$  be the outputs. (We would have  $m = 2^n$ ).
- ▶ Note that the subscripts here do not refer to different units.  
(That is why we are using different symbols).

- ▶ Define weighted geometric mean and weighted geometric mean of the complement as

$$G = \prod_i (O_i)^{P_i} \quad G' = \prod_i (1 - O_i)^{P_i}$$

- ▶ Define Normalized weighted geometric mean as

$$NWGM(O_1, \dots, O_m) = \frac{G}{G + G'}$$

- ▶ Then one can show that

$$NWGM(O_1, \dots, O_m) = \frac{1}{1 + ce^{-\beta \sum_i P_i S_i}} = f(ES_i)$$



- ▶ We can generalize this to a full network

$$y_j^l = f(\eta_j^l) = f\left(\sum_i w_{ij}^{l-1} \xi_i^{l-1} y_i^{l-1}\right)$$

we can show (under some approximation)

$$E[\eta_j^l] = \sum_i w_{ij}^{l-1} p_i^{l-1} E[y_i^{l-1}]$$

where  $p_j^l = E[\xi_j^l]$  ( $= p$ , normally).

- ▶ This is essentially the averaging effect that dropout provides.

# Batch Normalization

- ▶ In a neural network, we assume the input distribution to be constant. We learn a proper classifier for that distribution.
- ▶ As mentioned earlier, it helps to normalize the input distribution to have zero-mean and unit variance (or use a whitening transform)
- ▶ We can represent the output of a single hidden layer network as

$$y = F_2(F_1(x, \theta_1), \theta_2)$$

- ▶ Note that  $x' = F_1(x, \theta_1)$  is the output of hidden layer and is input to the final layer.
- ▶ But the distribution of  $x'$  keeps changing as we are learning  $\theta_1$ .

# Batch Normalization

- ▶ Batch normalization is a heuristic method that attempts to normalize the distribution of outputs at each layer in a deep network.
- ▶ Such normalization is seen to make learning more stable under SGD.
- ▶ This allows one to use larger step-sizes and hence achieve faster convergence.

- ▶ Suppose  $x$  is input to some layer (output of previous layer). The normalized version would be

$$\tilde{x} = \text{Normalize}(x, \mathcal{S})$$

where  $\mathcal{S}$  is the training set.

- ▶ But this defeats the idea of SGD because now gradient depends on all examples.
- ▶ Hence the normalization is done only on a minibatch.

- ▶ Let  $\mathcal{B} = \{x_1, \dots, x_m\}$  be the inputs to a layer in a minibatch.
- ▶ These are all vectors. But we would normalize each component separately to have mean 0 and variance 1. Hence, we show equations as if these are scalars.
- ▶ The normalization could be:

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m x_i \quad \sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2$$

$$\tilde{x}_i = \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

- ▶ The idea is to give  $\tilde{x}_i$  as input to next layer.

- ▶ In a deep network, each layer is transforming the representation.
- ▶ We do not know how the normalization would affect learning of the representation.
- ▶ Hence we compute the batch-normalization as

$$y_i = \text{BN}_{\gamma, \beta}(x_i) = \gamma \tilde{x}_i + \beta$$

and supply  $y_i$  as input to next layer.

- ▶ the parameters  $\gamma, \beta$  would also be learnt. (They would be different for different layers)
- ▶ If  $\gamma = 1, \beta = 0$  we are using usual normalization.
- ▶ If  $\gamma = \sqrt{\sigma_B^2 + \epsilon}, \beta = \mu_B$ , then there is no normalization.
- ▶ We would learn the appropriate normalization for each layer.

- ▶ In this notation, at some layer,  $x_i$  is the input to that layer in the original network which is now transformed into  $y_i$ .
- ▶ If  $L$  is the loss function, then we need  $\frac{\partial L}{\partial x_i}$  for updating weights connecting to this layer whose output is  $x_i$ .
- ▶ Since now  $y_i$  is input this layer, backpropagation can compute  $\frac{\partial L}{\partial y_i}$ .
- ▶ So, we need to express  $\frac{\partial L}{\partial x_i}$  in terms of  $\frac{\partial L}{\partial y_i}$  and other parameters of the batch-normalization transformation.
- ▶ We also need  $\frac{\partial L}{\partial \gamma}$  and  $\frac{\partial L}{\partial \beta}$ .
- ▶ All these can be obtained using chain rule of differentiation.

- ▶ Now, in the final learnt network there is a batch-normalization transformation after each layer:

$$y_i = \gamma \tilde{x}_i + \beta$$

- ▶ We are learning  $\gamma$  and  $\beta$  for each layer. But to compute  $\tilde{x}_i$  we need statistics of minibatch.
- ▶ What do we do during test time (regular operation)?



- ▶ After learning the network, with the learnt parameters fixed, we compute  $\mu_{\mathcal{B}}$  and  $\sigma_{\mathcal{B}}^2$  over many random minibatches from training set, all of size  $m$ .

- ▶ Let

$$\bar{\mu} = E[\mu_{\mathcal{B}}] \quad \text{and} \quad \bar{\sigma}_{\mathcal{B}}^2 = \frac{m}{m-1} E[\sigma_{\mathcal{B}}^2]$$

where  $E[\cdot]$  denotes average over all random minibatches.

- ▶ Let  $\gamma$  and  $\beta$  be learnt values for this layer. Then we use

$$y = \frac{\gamma}{\sqrt{\bar{\sigma}_{\mathcal{B}}^2 + \epsilon}} x + \left( \beta - \frac{\gamma \bar{\mu}}{\sqrt{\bar{\sigma}_{\mathcal{B}}^2 + \epsilon}} \right)$$

- ▶ This will be the final network that we use.

# Regularization for CNNs

- ▶ We have mentioned three methods of regularization.
- ▶ These are the main ones used.
- ▶ One may use any subset of them.
- ▶ These can be used with all feedforward networks (and also with recurrent networks)
- ▶ One also progressively decreases step-size as learning proceeds to promote more stable learning.

# Convolutional neural Networks

- ▶ CNNs are seen to achieve very high accuracies in a large number of applications involving classification of images, speech, text etc.
- ▶ The convolutional layer structure is very effective in extracting good feature representations using training data.
- ▶ CNNs can take a large part of the credit for the unprecedented interest in deep learning now.