


# Restricted Boltzmann Machine (RBM)

- ▶ Recurrent Neural Network with stochastic units
- ▶ A particular case of undirected graphical model.
- ▶ Belongs to what are often called energy-based models.
- ▶ A generative model (representing a probability distribution).
- ▶ Suitable for unsupervised learning.
- ▶ Can be used as a discriminative model too.


# Hopfield Network

- ▶ A fully connected recurrent neural network
- ▶ Binary units/neurons



# Hopfield Network

- ▶ Output (state) of  $i^{th}$  neuron:  $s_i \in \{0, 1\}$
- ▶  $w_{ij}$  – weight between  $i$  and  $j$   
We assume:  $w_{ij} = w_{ji}$ ,  $w_{ii} = 0$



# Hopfield Network – Dynamics

- ▶ Consider  $n$  neurons with outputs (states)  $s_1(k), \dots, s_n(k)$  at time  $k$ .
- ▶ At each  $k$ , only one randomly chosen neuron changes output.

$$s_i(k+1) = 1 \text{ iff } \sum_j w_{ij} s_j(k) + b_i > 0$$

- ▶ State of network:  $\mathbf{s}(k) = (s_1(k) \cdots s_n(k))$
- ▶  $\mathbf{s}(k)$  evolves over  $\{0, 1\}^n$ .

- ▶ Define energy,  $E : \{0, 1\}^n \rightarrow \Re$

$$E(\mathbf{s}) = -\frac{1}{2} \sum_{i,j} w_{ij} s_i s_j - \sum_i b_i s_i$$

- ▶ The dynamics of Hopfield network ensures that

$$E(\mathbf{s}(k+1)) - E(\mathbf{s}(k)) \leq 0, \forall k$$

- ▶ Dynamics converges to a minimum of energy.
- ▶ From current point we move to a 'neighbour' point in  $\{0, 1\}^n$ , seeking to minimize energy.

# Hopfield Network seeks minimum of energy

$$\begin{aligned} E(\mathbf{s}) &= -\frac{1}{2} \sum_{q,l} w_{ql} s_q s_l - \sum_q b_q s_q \\ &= -\frac{1}{2} \sum_{q,l \neq i} w_{ql} s_q s_l - s_i \sum_{q \neq i} w_{iq} s_q - \sum_{q \neq i} b_q s_q - b_i s_i \end{aligned}$$

- ▶ Suppose:  $s_j(k+1) = s_j(k)$ ,  $\forall j \neq i$ .
- ▶ Now we can write

$$E(\mathbf{s}(k+1)) - E(\mathbf{s}(k)) = (s_i(k) - s_i(k+1)) \left( \sum_{q \neq i} w_{iq} s_q(k) + b_i \right)$$

- ▶ The dynamics:  $s_i(k+1) = 1$ , iff  $\sum_{q \neq i} w_{iq} s_q(k) + b_i \geq 0$ .
- ▶ Hence  $E(\mathbf{s}(k+1)) - E(\mathbf{s}(k)) \leq 0$ .

- ▶ Hopfield network is one of the very influential early recurrent network models.
- ▶ Its dynamics seeks minima in an energy landscape
- ▶ Can be used as an optimizer.
- ▶ Can be used as associative memory (or for pattern completion) – a ‘generative’ model

# Boltzmann Machine

- ▶ Suppose we make the dynamics stochastic:  
at each instant, only one randomly selected neuron changes state; but the change is stochastic

$$\text{Prob}[s_i(k+1) = 1] = \frac{1}{1 + \exp(-\{\sum_j w_{ij}s_j(k) + b_i\})}$$

- ▶ Now,  $\mathbf{s}(k) = (s_1(k) \cdots s_n(k))$  would be a Markov chain over the state space  $\{0, 1\}^n$ .
- ▶ The transition structure is such that two states communicate iff they differ in only one coordinate.



- ▶ As before define energy by

$$E(\mathbf{s}) = -\frac{1}{2} \sum_{i,j} w_{ij} s_i s_j - \sum_i b_i s_i$$

- ▶ If  $\mathbf{s}, \mathbf{s}'$  differ only in  $i^{th}$  component, then

$$E(\mathbf{s}) - E(\mathbf{s}') = \left( \sum_{j \neq i} w_{ij} s_j + b_i \right) (s'_i - s_i)$$

- ▶ Let  $\mathbf{s}^1, \mathbf{s}^0 \in \{0, 1\}^n$  with  $s_j^1 = s_j^0, \forall j \neq i$  and  $s_i^1 = 1, s_i^0 = 0$ .
- ▶ Going from  $\mathbf{s}^0$  to  $\mathbf{s}^1$  involves choosing  $i^{th}$  neuron and setting it to 1. Hence this transition probability is

$$\begin{aligned}
 P(\mathbf{s}^0, \mathbf{s}^1) &= \frac{1}{n} \frac{1}{(1 + \exp(-\{\sum_j w_{ij} s_j^0(k) + b_i\}))} \\
 &= \frac{1}{n} \frac{1}{(1 + \exp((s_i^0 - s_i^1) \left( \{\sum_j w_{ij} s_j^0(k) + b_i\} \right)))} \\
 &= \frac{1}{n} \frac{1}{(1 + \exp(E(\mathbf{s}^1) - E(\mathbf{s}^0)))}
 \end{aligned}$$

- ▶ Going from  $\mathbf{s}^1$  to  $\mathbf{s}^0$  involves choosing  $i^{th}$  neuron and setting it to 0.
- ▶ Noting that  $1 - \frac{1}{1+e^{-x}} = \frac{1}{1+e^x}$

$$\begin{aligned}
 P(\mathbf{s}^1, \mathbf{s}^0) &= \frac{1}{n} \frac{1}{(1 + \exp(\{\sum_j w_{ij} s_j^0(k) + b_i\}))} \\
 &= \frac{1}{n} \frac{1}{(1 + \exp((s_i^1 - s_i^0) \left( \{\sum_j w_{ij} s_j^0(k) + b_i\} \right)))} \\
 &= \frac{1}{n} \frac{1}{(1 + \exp(E(\mathbf{s}^0) - E(\mathbf{s}^1)))}
 \end{aligned}$$

- ▶ Transitions are only to states differing in exactly one coordinate
- ▶ The transition structure is

$$P(\mathbf{s}, \mathbf{s}') = \frac{1}{n} \frac{1}{(1 + \exp(E(\mathbf{s}') - E(\mathbf{s})))}$$

- ▶ The chain prefers going to lower energy states.
- ▶ This Markov chain is finite, irreducible, aperiodic and hence ergodic.
- ▶ The stationary distribution is given by

$$p(\mathbf{s}) = \frac{1}{Z} \exp(-E(\mathbf{s}))$$

# The stationary distribution of this chain

- ▶ if  $p$  satisfies

$$p(s')P(s', s) = p(s)P(s, s'), \quad \text{for all states } s, s'$$

then it is the stationary distribution

- ▶ We can easily see that

$$\frac{\exp(-E(s'))}{1 + \exp(E(s) - E(s'))} = \frac{\exp(-E(s))}{1 + \exp(E(s') - E(s))}, \quad \forall s, s'$$

- ▶ Hence, the stationary distribution is

$$p(s) = \frac{1}{Z} \exp(-E(s))$$

where  $Z = \sum_s \exp(-E(s))$  is the normalizing constant.

# Energy based Model

- ▶ The boltzman machine is a generative model.
- ▶ It represents a distribution over  $\{0, 1\}^n$  given by

$$p(\mathbf{s}) = \frac{1}{Z} \exp(-E(\mathbf{s}))$$

where  $Z$  is the normalizing constant and

$$E(\mathbf{s}) = -\frac{1}{2} \sum_{ij} w_{ij} s_i s_j - \sum_i b_i s_i$$

- ▶ Each configuration (point in  $\{0, 1\}^n$ ) has an energy which determines its probability.

- ▶ Boltzmann machine can also be used as a discriminative model.
- ▶ If we hold some neurons at some values, then at steady state what we see at the other neurons is the conditional distribution.
- ▶ This is how we can use it as classifier.

- ▶ We have a generative model:  $p(\mathbf{s})$ .
- ▶ We learn this probability distribution from given data.
- ▶ We can use it as a discriminative model by designating some components of  $\mathbf{s}$  as 'input' and some as 'output'
- ▶ However, we learn the probability distribution through Unsupervised learning – 'treat all variables uniformly'



# Need for hidden units

- ▶ We learn a generative model:  $p(\mathbf{s})$ :

$$p(\mathbf{s}) = \frac{1}{Z} \exp(-E(\mathbf{s}))$$

where  $Z$  is the normalizing constant and

$$E(\mathbf{s}) = -\frac{1}{2} \sum_{ij} w_{ij} s_i s_j - \sum_i b_i s_i$$

- ▶ The probability depends on energy which is quadratic
- ▶ We can essentially model only upto second-order statistics.
- ▶ Hence we need hidden nodes to learn 'proper representation'

# Need for hidden units

- ▶ Let  $n = 3$ .
- ▶ Let distribution  $p_1$ : uniform over  $\{0, 1\}^3$ .
- ▶ Let distribution  $p_2$ : uniform only over the four tuples:  $(0, 0, 0), (0, 1, 1), (1, 0, 1), (1, 1, 0)$ . (The other four tuples have zero probability)
- ▶ But  $p_1$  and  $p_2$  have identical statistics upto second order. ( $P[s_i s_j = 1], P[s_i = 1]$  are same under both)
- ▶ We can not distinguish between  $p_1$  and  $p_2$  with our current formulation
- ▶ Note that  $p_2$  models (or represents) the XOR function

# Boltzmann Machine

- ▶ We would have  $(m + n)$  units:  $v_1, \dots, v_m, h_1, \dots, h_n$ .
- ▶ The  $v_i$  are called visible and the  $h_i$  are called hidden.
- ▶ The formulation is still the same. The network is fully connected. (Simply some of the  $s_i$  are visible while others are hidden).
- ▶ The Energy is defined the same way and the machine represents a distribution

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \exp(-E(\mathbf{v}, \mathbf{h}))$$

- ▶ But we are only interested in the distribution on visible units which is given by

$$p(\mathbf{v}) = \sum_{\mathbf{h}} \frac{1}{Z} \exp(-E(\mathbf{v}, \mathbf{h}))$$

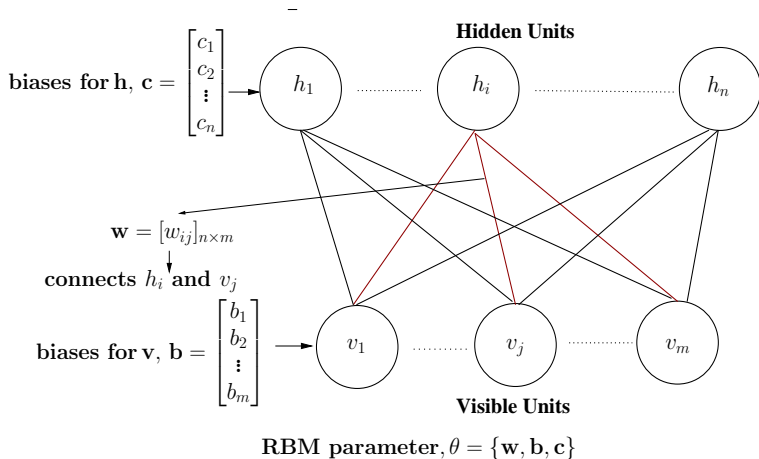
- ▶ Training data would be samples of visible units only
- ▶ The hidden units are there to help us realize the needed distribution on the visible units.
- ▶ This is the original Boltzmann machine proposed.

- ▶ The weights in Boltzmann machine are learnt (or estimated) through maximum likelihood estimation.
- ▶ But training is slow and computationally intensive
- ▶ A conditional independence assumption makes the model more tractable.
- ▶ That is the Restricted Boltzmann Machine

# Restricted Boltmann Machine

- ▶ As earlier we have  $m$  visible and  $n$  hidden units:  
 $v_1, \dots, v_m, h_1, \dots, h_n$ .
- ▶ But connections are only between visible and hidden units;  
no connections from visible to visible or hidden to hidden.
- ▶ Notation:  
 $w_{ij}$  is weight of connection between  $v_j$  and  $h_i$   
 $b_i$  is bias for  $v_i$   
 $c_i$  is bias for  $h_i$

# RBM



- ▶ Now the energy is given by

$$\begin{aligned} E(\mathbf{v}, \mathbf{h}) &= - \sum_{i=1}^n \sum_{j=1}^m w_{ij} h_i v_j - \sum_{j=1}^m b_j v_j - \sum_{i=1}^n c_i h_i \\ &= -\mathbf{h}^T W \mathbf{v} - \mathbf{b}^T \mathbf{v} - \mathbf{c}^T \mathbf{h} \end{aligned}$$

- ▶ The probability distribution, over  $\{0, 1\}^{(m+n)}$ , represented by the RBM is

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \exp(-E(\mathbf{v}, \mathbf{h}))$$



# Conditional Independence

- Consistent with the interconnections, we get the following conditional independence:

$$p(\mathbf{h}|\mathbf{v}) = \prod_i p(h_i|\mathbf{v}); \quad p(\mathbf{v}|\mathbf{h}) = \prod_j p(v_j|\mathbf{h})$$

- The conditional probabilities are given by

$$\begin{aligned} \text{Prob}[h_i = 1|\mathbf{v}] &= \text{sig} \left( \sum_{j=1}^m w_{ij} v_j + c_i \right) \\ \text{Prob}[v_i = 1|\mathbf{h}] &= \text{sig} \left( \sum_{j=1}^n w_{ji} h_j + b_i \right) \end{aligned}$$

where  $\text{sig}(x) = 1/(1 + \exp(-x))$ .

- We have

$$\begin{aligned}\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} &= \sum_{\mathbf{h}} e^{\mathbf{b}^T \mathbf{v}} e^{\sum_{i=1}^n h_i (c_i + \sum_{j=1}^m w_{ij} v_j)} \\ &= e^{\mathbf{b}^T \mathbf{v}} \sum_{\mathbf{h}} \prod_{i=1}^n e^{h_i (c_i + \sum_{j=1}^m w_{ij} v_j)} \\ &= e^{\mathbf{b}^T \mathbf{v}} \sum_{h_1} e^{h_1 (c_1 + \sum_{j=1}^m w_{1j} v_j)} \dots \sum_{h_n} e^{h_n (c_n + \sum_{j=1}^m w_{nj} v_j)} \\ &= e^{\mathbf{b}^T \mathbf{v}} \prod_{i=1}^n (1 + e^{c_i + \sum_{j=1}^m w_{ij} v_j})\end{aligned}$$

- This gives us an expression for the marginal distribution on the visible units

$$\begin{aligned} p(\mathbf{v}) &= \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \\ &= \frac{1}{Z} e^{\mathbf{b}^T \mathbf{v}} \prod_{i=1}^n (1 + e^{c_i + \sum_{j=1}^m w_{ij} v_j}) \end{aligned}$$

- ▶ This also gives us the conditional distribution

$$\begin{aligned} p(\mathbf{h}|\mathbf{v}) &= \frac{e^{-E(\mathbf{v},\mathbf{h})}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v},\mathbf{h})}} \\ &= \frac{e^{\mathbf{b}^T \mathbf{v}} e^{\sum_{i=1}^n h_i (c_i + \sum_{j=1}^m w_{ij} v_j)}}{e^{\mathbf{b}^T \mathbf{v}} \prod_{i=1}^n (1 + e^{c_i + \sum_{j=1}^m w_{ij} v_j})} \\ &= \prod_{i=1}^n \frac{e^{h_i (c_i + \sum_{j=1}^m w_{ij} v_j)}}{(1 + e^{c_i + \sum_{j=1}^m w_{ij} v_j})} \end{aligned}$$

- ▶ Shows that  $h_i$  are conditionally independent given  $\mathbf{v}$

- Now we get

$$\begin{aligned} p(h_i|\mathbf{v}) &= \sum_{h_j, j \neq i} \prod_{j=1}^n \frac{e^{h_j(c_j + \sum_{q=1}^m w_{jq}v_q)}}{(1 + e^{c_j + \sum_{q=1}^m w_{jq}v_q})} \\ &= \frac{e^{h_i(c_i + \sum_{q=1}^m w_{iq}v_q)}}{(1 + e^{c_i + \sum_{q=1}^m w_{iq}v_q})} \end{aligned}$$

- This gives us

$$\text{Prob}[h_i = 1|\mathbf{v}] = \text{sig} \left( c_i + \sum_{j=1}^m w_{ij}v_j \right)$$

- A similar expression holds for  $\text{Prob}[v_i = 1|\mathbf{h}]$

# Summary of RBM model

- ▶ The RBM is a generative model with  $m$  visible and  $n$  hidden nodes:

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} \exp(-E(\mathbf{v}, \mathbf{h}))$$

where the energy is given by

$$E(\mathbf{v}, \mathbf{h}) = -\mathbf{h}^T W \mathbf{v} - \mathbf{b}^T \mathbf{v} - \mathbf{c}^T \mathbf{h}$$

and the normalizing constant,  $Z$ , is given by

$$Z = \sum_{\mathbf{v}, \mathbf{h}} \exp(-E(\mathbf{v}, \mathbf{h}))$$

## Summary . . .

- ▶ The following conditional independences hold:


$$p(\mathbf{h}|\mathbf{v}) = \prod_i p(h_i|\mathbf{v}); \quad p(\mathbf{v}|\mathbf{h}) = \prod_j p(v_j|\mathbf{h})$$

- ▶ The conditional probabilities are

$$\begin{aligned} \text{Prob}[h_i = 1|\mathbf{v}] &= \text{sig} \left( \sum_{j=1}^m w_{ij} v_j + c_i \right) \\ \text{Prob}[v_i = 1|\mathbf{h}] &= \text{sig} \left( \sum_{j=1}^n w_{ji} h_j + b_i \right) \end{aligned}$$

# RBM as graphical model

- RBM can be specified as an undirected graphical model.



- Any two hidden nodes are conditionally independent given all visible nodes and vice versa



## Recall: Undirected graphical models

- ▶ Suppose a generative model for  $\mathbf{X} = (X_1, \dots, X_n)$  is an undirected graphical model. Let  $\mathcal{C}$  be the set of maximal cliques.
- ▶ Then the joint distribution should be of the form

$$p(\mathbf{x}) = \frac{1}{Z} e^{\sum_c \psi_c(\mathbf{x})}$$

where the  $\psi_c$  are called clique potentials.

- ▶  $\psi_c(\mathbf{x})$  depends only on those components of  $\mathbf{x}$  that are in  $c$ .  
(Hammersley-Clifford Theorem)

# Case of RBM

- ▶ Thus every MRF is an energy-based model with energy function given by

$$E(\mathbf{x}) = - \sum_{c \in \mathcal{C}} \psi_c(\mathbf{x})$$



- ▶ In an RBM, the cliques are  $\{v_i, h_j\}$ . Since these are binary, a general energy function would have terms involving  $v_i$ ,  $h_j$  and  $v_i h_j$
- ▶ The graphical model comes from conditional independence we imposed
- ▶ For the given graphical model, our choice of energy function is very general.

# Representational Power of RBMs

- ▶ As said earlier, we are only interested in modelling the distribution over visible units.
- ▶ The hidden units are there only so that we can represent the desired distribution.
- ▶ Question: What distributions over  $\{0, 1\}^m$  are representable by an RBM with  $m$  visible nodes (given sufficient number of hidden nodes)
- ▶ Given  $n = 2^m + 1$  hidden nodes, RBM can represent any distribution over  $\{0, 1\}^m$ .
- ▶ If the target distribution is supported on fewer terms, then we can do with lesser number of hidden nodes.

# RBM as a discriminative model

- ▶ We can use RBM as a classifier also.
- ▶ We can take  $\mathbf{v} \in \{0, 1\}^m$  as 'feature vector' and take  $\mathbf{y} \in \{0, 1\}^K$  as class label. (Note  $\mathbf{y}$  takes only  $K$  distinct values in a  $K$ -class problem)
- ▶ For RBM,  $(\mathbf{v}, \mathbf{y})$  together make all the visible units and we can have  $n$  hidden units.



- ▶ Now the training data would be labelled samples.
- ▶ We would learn the distribution of  $(\mathbf{v}, \mathbf{y})$

- ▶ Let
  - $w_{ij}$  be weight between  $h_i$  and  $v_j$
  - $\tilde{w}_{ij}$  be weight between  $y_j$  and  $h_i$ .
  - $b_i$  be bias for  $v_i$ ,  $\tilde{b}_i$  for  $y_i$  and  $c_i$  for  $h_i$
- ▶ Now we can write the probability model of RBM as

$$p(\mathbf{v}, \mathbf{y}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{y}, \mathbf{h})}$$

where the energy is given by

$$E(\mathbf{v}, \mathbf{y}, \mathbf{h}) = -\mathbf{h}^T W \mathbf{v} - \mathbf{b}^T \mathbf{v} - \mathbf{c}^T \mathbf{h} - \mathbf{h}^T \tilde{W} \mathbf{y} - \tilde{\mathbf{b}}^T \mathbf{y}$$

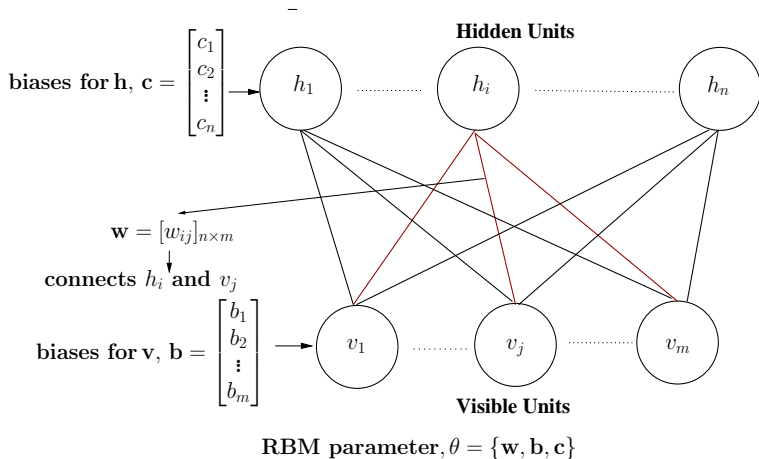
- ▶ As mentioned earlier  $\mathbf{y}$  takes only the  $K$  unit vectors as values.
- ▶ After learning all the weights we need to predict  $\mathbf{y}$  given  $\mathbf{v}$ .
- ▶ We calculate

$$p(\mathbf{y}|\mathbf{v}) = \frac{\sum_{\mathbf{h}} p(\mathbf{v}, \mathbf{y}, \mathbf{h})}{\sum_{\mathbf{h}, \mathbf{y}} p(\mathbf{v}, \mathbf{y}, \mathbf{h})}$$

for all the unit vectors and take the one with highest probability.



# Learning RBM



Need to learn parameters from training data.

- ▶ Given some value for parameters,  $\theta$ , RBM represents a probability density  $p_{\theta}(\mathbf{v})$ .
- ▶ Given *iid* training data from some distribution, we want to learn a parametric model of the distribution.
- ▶ Hence we can use ML estimation of  $\theta$ .

- ▶ The training data:  $\mathcal{D} = \{\mathbf{v}^1, \mathbf{v}^2, \dots, \mathbf{v}^M\}$
- ▶ The data likelihood is

$$\mathcal{L}(\theta|\mathcal{D}) = \prod_{k=1}^M p_{\theta}(\mathbf{v}^k)$$

- ▶ The log-likelihood is given by

$$\ell(\theta) = \ln(\mathcal{L}(\theta)) = \sum_{k=1}^M \ln(p_{\theta}(\mathbf{v}^k))$$

- ▶ We can minimize the negative log-likelihood (NLL) to get the ML estimate.
- ▶ (Recall that) Minimizing NLL is same as minimizing KL divergence between  $p_{\text{data}}$  and  $p_{\theta}$

- ▶ For simplicity of notation, assume we have only one training sample,  $\mathbf{v}^1$ .
- ▶ The log-likelihood is

$$\ell(\theta) = \ln(p_{\theta}(\mathbf{v}^1)) = \ln \left( \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})} \right)$$

- ▶ We want to do gradient descent on negative log-likelihood.

$$\frac{\partial \ell(\theta)}{\partial \theta} = \frac{\partial}{\partial \theta} \left( \ln \left( \sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})} \right) - \ln(Z) \right)$$

- ▶ we have

$$\frac{\partial \ell(\theta)}{\partial \theta} = \frac{\partial}{\partial \theta} \left( \ln \left( \sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})} \right) - \ln(Z) \right)$$

- ▶ The parameter vector,  $\theta$ , consists of all weights,  $w_{ij}$  and the biases  $b_i$  and  $c_i$ .
- ▶ The above is a vector equation with one derivative for each component of parameter vector.
- ▶ Let us write  $\frac{\partial}{\partial w}$  for any one component. ( $w$  may be  $w_{ij}$ ,  $b_i$  or  $c_i$ ).

- Noting that  $Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}$ , we have

$$\begin{aligned} \frac{\partial \ell(\theta)}{\partial w} &= \frac{\partial}{\partial w} \ln \left( \sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})} \right) - \frac{\partial}{\partial w} \ln(Z) \\ &= \frac{\sum_{\mathbf{h}} \frac{\partial e^{-E(\mathbf{v}^1, \mathbf{h})}}{\partial w}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})}} - \frac{\sum_{\mathbf{v}, \mathbf{h}} \frac{\partial e^{-E(\mathbf{v}, \mathbf{h})}}{\partial w}}{Z} \\ &= f(\theta) - g(\theta), \text{ (say)} \end{aligned}$$

- The first term can be simplified as follows

$$\begin{aligned} f(\theta) &= \sum_{\mathbf{h}} \frac{\frac{\partial e^{-E(\mathbf{v}^1, \mathbf{h})}}{\partial w}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})}} \\ &= \sum_{\mathbf{h}} \left( \frac{e^{-E(\mathbf{v}^1, \mathbf{h})}}{\sum_{\mathbf{h}} e^{-E(\mathbf{v}^1, \mathbf{h})}} \frac{\partial(-E(\mathbf{v}^1, \mathbf{h}))}{\partial w} \right) \\ &= \sum_{\mathbf{h}} p(\mathbf{h}|\mathbf{v}^1) \frac{\partial(-E(\mathbf{v}^1, \mathbf{h}))}{\partial w} \\ &= \mathcal{E}_{\mathbf{h}|\mathbf{v}^1} \left[ \frac{\partial(-E(\mathbf{v}^1, \mathbf{h}))}{\partial w} \right] \end{aligned}$$

where  $\mathcal{E}_{\mathbf{h}|\mathbf{v}^1}$  denotes expectation with respect to the conditional distribution of  $\mathbf{h}$  given  $\mathbf{v}$  at  $\mathbf{v} = \mathbf{v}^1$  (and at the current values of parameters).

- Now let us look at the second term in gradient of log-likelihood

$$\begin{aligned} g(\theta) &= \frac{\sum_{\mathbf{v}, \mathbf{h}} \frac{\partial e^{-E(\mathbf{v}, \mathbf{h})}}{\partial w}}{Z} \\ &= \sum_{\mathbf{v}, \mathbf{h}} \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{Z} \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \\ &= \sum_{\mathbf{v}, \mathbf{h}} p(\mathbf{v}, \mathbf{h}) \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \\ &= \mathcal{E}_{\mathbf{v}, \mathbf{h}} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] \end{aligned}$$

where  $\mathcal{E}_{\mathbf{v}, \mathbf{h}}$  denotes expectation with respect to the joint distribution of  $(\mathbf{v}, \mathbf{h})$  (at the current parameter values,  $\theta$ ).



- ▶ In the above we have taken only one example and set  $\ell(\theta) = \ln(p(\mathbf{v}^1))$ .
- ▶ But we have  $M$  samples and hence have to take sum over them for log-likelihood.
- ▶ We need to maximize

$$\ell(\theta) = \frac{1}{M} \sum_{k=1}^M \ln(p(\mathbf{v}^k))$$

- ▶ We have taken average log-likelihood above which does not affect the optimization but results in interesting interpretation for the final algorithm.

- Now, the earlier function  $f(\theta)$  becomes

$$\begin{aligned} f(\theta) &= \frac{1}{M} \sum_{k=1}^M \mathcal{E}_{\mathbf{h}|\mathbf{v}^k} \left[ \frac{\partial(-E(\mathbf{v}^k, \mathbf{h}))}{\partial w} \right] \\ &= \sum_{k=1}^M p_{\text{data}}(\mathbf{v}^k) \mathcal{E}_{\mathbf{h}|\mathbf{v}^k} \left[ \frac{\partial(-E(\mathbf{v}^k, \mathbf{h}))}{\partial w} \right] \\ &= \mathcal{E}_{\text{data}} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] \end{aligned}$$

where  $\mathcal{E}_{\text{data}}$  denotes expectation w.r.t. 'distribution given by data'.

- The second term  $g(\theta)$  now becomes

$$\begin{aligned} g(\theta) &= \frac{1}{M} \sum_{k=1}^M \mathcal{E}_{\mathbf{v}, \mathbf{h}} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] \\ &= \mathcal{E}_{\mathbf{v}, \mathbf{h}} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] \\ &= \mathcal{E}_{\theta} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] \end{aligned}$$

where  $\mathcal{E}_{\theta}$  represents expectation with respect to the distribution 'represented by the parameter vector  $\theta$ '.

- ▶ Thus, finally we get the gradient of log-likelihood is

$$\frac{\partial \ell(\theta)}{\partial w} = \mathcal{E}_{\text{data}} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] - \mathcal{E}_{\theta} \left[ \frac{\partial(-E(\mathbf{v}, \mathbf{h}))}{\partial w} \right]$$

- ▶ Or, the gradient of negative log-likelihood is

$$\frac{\partial(-\ell(\theta))}{\partial w} = \mathcal{E}_{\text{data}} \left[ \frac{\partial(E(\mathbf{v}, \mathbf{h}))}{\partial w} \right] - \mathcal{E}_{\theta} \left[ \frac{\partial(E(\mathbf{v}, \mathbf{h}))}{\partial w} \right]$$

- ▶ Note that the above is true for general Boltzmann machine as well because we have not used the conditional independences anywhere.

- By definition of energy, we have

$$\frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial w_{ij}} = -h_i v_j; \quad \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial b_i} = -v_i; \quad \frac{\partial E(\mathbf{v}, \mathbf{h})}{\partial c_i} = -h_i$$

- Hence we get, for gradient descent on negative log-likelihood

$$\Delta w_{ij} = -\frac{\partial(-\ell(\theta))}{\partial w_{ij}} = (\mathcal{E}_{\text{data}}[h_i v_j] - \mathcal{E}_{\theta}[h_i v_j])$$

- Similar updating is done for  $b_i$  and  $c_i$ .
- This is the learning algorithm for the original Boltzmann machine too.

- ▶ Let us look at how we can calculate the first term

$$\begin{aligned}\mathcal{E}_{\text{data}}[h_i v_j] &= \sum_{k=1}^M (1/M) v_j^k \text{Prob}[h_i = 1 | \mathbf{v}^k] \\ &= \sum_{k=1}^M (1/M) v_j^k \text{sig} \left( \sum_{s=1}^m w_{si} v_s^k + c_i \right)\end{aligned}$$

- ▶ Note that we have used the RBM conditional distribution here.
- ▶ The  $(1/M)$  factor is not important. For each training sample,  $w_{ij}$  is updated by  $v_j^k \text{Prob}[h_i = 1 | \mathbf{v}^k]$

- ▶ Now let us look at the second term in the update equation.

$$\mathcal{E}_\theta [h_i v_j] = \sum_{\mathbf{v}, \mathbf{h}} p(\mathbf{v}, \mathbf{h}) h_i v_j$$

- ▶ Computationally hard to – summation is over  $2^{m+n}$  terms.
- ▶ We can try the same decomposition of the expectation as in the first term

$$\begin{aligned}\mathcal{E}_\theta [h_i v_j] &= \mathcal{E}_{p(\mathbf{v})} \mathcal{E}_{p(\mathbf{h}|\mathbf{v})} [h_i v_j] \\ &= \sum_{\mathbf{v}} p(\mathbf{v}) v_j \text{Prob}[h_i = 1|\mathbf{v}]\end{aligned}$$

- ▶ Still hard to compute – summation is over  $2^m$  terms.
- ▶ Hence we approximate the expectation by sample mean.

- ▶ The gradient of the log-likelihood contains two terms – both are expectations.
- ▶ One of them is easy to compute for the RBM because of the conditional independence.
- ▶ But the second term is intractable and hence expectation is approximated by sample mean.
- ▶ For the general Boltzmann machine, both terms are intractable and both terms need approximation through sample mean.



- ▶ Let us consider an algorithm where we use a single sample.
- ▶ Let  $\mathbf{v}'$  a training sample. Let  $\mathbf{v}$  be a sample from the distribution determined by current  $\theta$ . Then

$$\Delta w_{ij} = v'_j \text{Prob}[h_i = 1 | \mathbf{v}'] - v_j \text{Prob}[h_i = 1 | \mathbf{v}]$$

- ▶ All we need is a method to sample from the distribution  $p(\mathbf{v})$  (at the current parameter values).

# CD( $k$ ) Algorithm

- ▶ The distribution we are interested in is the stationary distribution of the Markov chain.
- ▶ So, we essentially run the chain for  $k$  steps and take the state as the sample.
- ▶ Gradient descent on NLL with this approximate gradient is called the Contrastive Divergence algorithm.
- ▶ To obtain the sample, we run the chain through what is called Gibbs sampling.
- ▶ We first look at Hastings-Metropolis algorithm and explain Gibbs sampling through that.

# Hastings-Metropolis Sampling

- ▶ Let  $\mathbf{X}$  be a random vector with mass function

$$p(\mathbf{x}) = b(\mathbf{x})/Z$$

We want samples from this without knowing  $Z$ .

- ▶ Let  $q(\mathbf{x}, \mathbf{x}')$  (called the proposal distribution) be the transition probabilities of some ergodic chain on the same state space.
- ▶ Define a Markov chain with transition probabilities

$$P(\mathbf{x}, \mathbf{x}') = q(\mathbf{x}, \mathbf{x}')\alpha(\mathbf{x}, \mathbf{x}'), \quad \mathbf{x} \neq \mathbf{x}'$$

At  $\mathbf{x}$ , generate  $\mathbf{x}'$  using  $q$  and accept it with probability  $\alpha(\mathbf{x}, \mathbf{x}')$  (stay at  $\mathbf{x}$  with remaining prob)

- ▶ Given a  $q$ , we choose  $\alpha$  as

$$\alpha(\mathbf{x}, \mathbf{x}') = \min \left( \frac{p(\mathbf{x}')q(\mathbf{x}', \mathbf{x})}{p(\mathbf{x})q(\mathbf{x}, \mathbf{x}')}, 1 \right) = \min \left( \frac{b(\mathbf{x}')q(\mathbf{x}', \mathbf{x})}{b(\mathbf{x})q(\mathbf{x}, \mathbf{x}')}, 1 \right)$$

- ▶ Note that one of  $\alpha(\mathbf{x}, \mathbf{x}')$ ,  $\alpha(\mathbf{x}', \mathbf{x})$  has to be 1.
- ▶ Then one can show that the Markov chain (corresponding to  $P$ ) has  $p(\cdot)$  as the stationary distribution.

- Suppose

$$\alpha(\mathbf{x}, \mathbf{x}') = \frac{p(\mathbf{x}')q(\mathbf{x}', \mathbf{x})}{p(\mathbf{x})q(\mathbf{x}, \mathbf{x}')} < 1 \quad (\text{and hence } \alpha(\mathbf{x}', \mathbf{x}) = 1)$$

- Then we have

$$\begin{aligned} p(\mathbf{x})P(\mathbf{x}, \mathbf{x}') &= p(\mathbf{x})q(\mathbf{x}, \mathbf{x}')\alpha(\mathbf{x}, \mathbf{x}') \\ &= p(\mathbf{x}')q(\mathbf{x}', \mathbf{x}) \\ &= p(\mathbf{x}')q(\mathbf{x}', \mathbf{x})\alpha(\mathbf{x}', \mathbf{x}) \\ &= p(\mathbf{x}')P(\mathbf{x}', \mathbf{x}) \end{aligned}$$

thus showing  $p(\mathbf{x})$  is the stationary distribution.

# Gibbs Sampling

- ▶ Gibbs sampling corresponds to a special case of  $q$ .
- ▶ Let  $\mathbf{x} = (x_1, \dots, x_m)$ . Let  $\mathbf{x}_{-i}$  denote a vector containing all components of  $\mathbf{x}$  except the  $i^{th}$  one.
- ▶ In Gibbs sampling, we generate successive states as

$$(x_1, x_2, \dots, x_m) \rightarrow (x'_1, x_2, \dots, x_m) \rightarrow (x'_1, x'_2, \dots, x_m) \cdots$$

where  $x'_i$  is generated from the distribution  $p(x_i | \mathbf{x}_{-i})$

- ▶ This corresponds to a  $q$  with  $q(\mathbf{x}, \mathbf{x}') = 0$  if  $\mathbf{x}, \mathbf{x}'$  differ in more than one component and is given by the above conditional distribution otherwise.

- If we use this  $q$  in Hastings-Metropolis, then  $\alpha$  would always be 1.

$$\begin{aligned}\alpha(\mathbf{x}, \mathbf{x}') &= \frac{p(\mathbf{x}')q(\mathbf{x}', \mathbf{x})}{p(\mathbf{x})q(\mathbf{x}, \mathbf{x}')} \quad (\mathbf{x}, \mathbf{x}' \text{ differ in } i^{th}) \\ &= \frac{p(x'_i | \mathbf{x}'_{-i})p(\mathbf{x}'_{-i})p(x_i | \mathbf{x}'_{-i})}{p(x_i | \mathbf{x}_{-i})p(\mathbf{x}_{-i})p(x'_i | \mathbf{x}_{-i})} \\ &= \frac{p(x'_i | \mathbf{x}_{-i})p(\mathbf{x}_{-i})p(x_i | \mathbf{x}_{-i})}{p(x_i | \mathbf{x}_{-i})p(\mathbf{x}_{-i})p(x'_i | \mathbf{x}_{-i})} = 1\end{aligned}$$

because  $\mathbf{x}_{-i} = \mathbf{x}'_{-i}$ .

# Gibbs Chain for RBM

- ▶ Due to the conditional independences in the RBM graph, Gibbs sampling is efficient here.
- ▶ For any  $h_i$  the conditional distribution given all the other nodes is same as that given only the visible nodes.
- ▶ Hence, given current  $\mathbf{v}$  we can choose the next values for all  $h_i$  together using the corresponding conditional distributions.
- ▶ Similarly, given current  $\mathbf{h}$  we can choose the next values for all  $v_i$  together.
- ▶ This is how the Gibbs chain for the RBM is run.



# CD( $k$ ) Algorithm

- ▶ We can now summarize the CD( $k$ ) algorithm.
- ▶ Let  $\mathbf{v}^{(0)}$  denote a training sample.
- ▶ We initialize the chain with this and run it for  $k$  steps.  
The state of visible units be  $\mathbf{v}^{(k)}$
- ▶ We update  $w_{ij}$  as

$$\Delta w_{ij} = v_j^{(0)} \text{Prob}[h_i = 1 | \mathbf{v}^{(0)}] - v_j^{(k)} \text{Prob}[h_i = 1 | \mathbf{v}^{(k)}]$$

(Similar updates are done for  $b_i$  and  $c_i$ )

- ▶ We need to do this with each training sample.

- ▶ In  $CD(k)$  we are approximating the expectation with a single sample.
- ▶ Only as  $k$  tends to infinity, the sample would be from the needed distribution. For finite  $k$  it is a biased estimate.
- ▶  $CD(k)$  is quite effective in practice.

# A Variation – PCD

- ▶ A variation is Persistent CD or PCD.
- ▶ Here, instead of initializing the Gibbs chain with the training sample each time, we let the Gibbs chain run independently.
- ▶ After each parameter update, the Gibbs chain is started in the last state of the previous Gibbs chain.
- ▶ The idea is that this results in a better ‘mixing’
- ▶ PCD results in better learning in terms of smoother and better increase in log-likelihood.
- ▶ There is also a variation that maintains multiple Gibbs chains called Parallel Tempering.



- ▶ There are many generalizations possible on RBMs
- ▶ RBMs are good for unsupervised learning of representations.
- ▶ The representation on the hidden layer can be a good feature representation.
- ▶ Hence one variation is the so called convolutional RBMs.
- ▶ Here, we can have local connectivity (and also weight sharing with multiple feature planes)
- ▶ Found useful for feature detection in, e.g., speech

- ▶ So far we considered both  $v_i$  and  $h_j$  to be binary.
- ▶ Training data is for visible units.
- ▶ So, this RBM can learn distribution for binary data only
- ▶ We may want to learn distributions over  $\mathbb{R}^n$ .
- ▶ We want RBMs where output of visible unit is a real number.
- ▶ Such an extension is possible.

# Gaussian-Binary RBM

- ▶ We take hidden units to be still binary.
- ▶ But we take visible units to be real-valued.
- ▶ We make the same conditional independence assumptions. (Same graphical model)
- ▶ We retain the same conditional distribution of  $h_i$  given  $\mathbf{v}$ .
- ▶ We can get a model where the conditional distribution on visible units given  $\mathbf{h}$  is Gaussian.

# Gaussian-Binary RBM

- ▶ We take  $\mathbf{v} \in \Re^m$  and  $\mathbf{h} \in \{0, 1\}^n$ .
- ▶ We can prescribe the conditional distributions as below:

$$p(\mathbf{v}|\mathbf{h}) \sim \mathcal{N}(\mathbf{v}; \mathbf{b} + W\mathbf{h}, \sigma^2 I)$$

$$\text{Prob}[h_i = 1|\mathbf{v}] = \text{sig} \left( c_i + \frac{1}{\sigma^2} \sum_{j=1}^m w_{ij} v_j \right)$$

- ▶ We can define energy as below to realize the above distributions

$$E(\mathbf{v}, \mathbf{h}) = -\frac{1}{\sigma^2} \mathbf{v}^T W \mathbf{h} + \frac{1}{2\sigma^2} \|\mathbf{v} - \mathbf{b}\|^2 - \mathbf{c}^T \mathbf{h}$$

- ▶ The only extra parameter here is  $\sigma^2$  which is the variance of the Gaussian.