

Recap

- ▶ We are discussing the SVM method for learning classifiers.
- ▶ Objective is to learn the optimal hyperplane – one that maximizes the ‘margin of separation’.
- ▶ It can be formulated as a constrained optimization problem.

Linear SVM – data linearly separable

- ▶ The optimal hyperplane is a solution of

$$\begin{array}{ll}\text{minimize} & \frac{1}{2}W^TW \\ \text{subject to} & y_i(W^TX_i + b) \geq 1, \quad i = 1, \dots, n\end{array}$$

- ▶ We solve the dual given by

$$\begin{array}{ll}\max_{\boldsymbol{\mu}} & q(\boldsymbol{\mu}) = \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j X_i^T X_j \\ \text{subject to} & \mu_i \geq 0, \quad i = 1, \dots, n, \quad \sum_{i=1}^n y_i \mu_i = 0\end{array}$$

- ▶ Then the final solution is:

$$W^* = \sum \mu_i^* y_i X_i, \quad b^* = y_j - X_j^T W^*, \quad j \text{ such that } \mu_j > 0$$

Linear SVM

- ▶ The primal problem is

$$\begin{aligned} &\text{minimize} && \frac{1}{2}W^TW + C \sum_{i=1}^n \xi_i \\ &\text{subject to} && y_i(W^TX_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ &&& \xi_i \geq 0, \quad i = 1, \dots, n \end{aligned}$$

- ▶ The dual problem is:

$$\begin{aligned} &\max_{\boldsymbol{\mu}} && q(\boldsymbol{\mu}) = \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j X_i^T X_j \\ &\text{subject to} && 0 \leq \mu_i \leq C, \quad i = 1, \dots, n, \quad \sum_{i=1}^n y_i \mu_i = 0 \end{aligned}$$

- ▶ We solve dual and the final optimal hyperplane is

$$\begin{aligned} W^* &= \sum \mu_i^* y_i X_i, \\ b^* &= y_j - X_j^T W^*, \quad (j \text{ s.t. } 0 < \mu_j < C). \end{aligned}$$

The Linear SVM

- ▶ Given training data, $(X_i, y_i), i = 1, \dots, n$, we solve

$$\max_{\boldsymbol{\mu}} \quad q(\boldsymbol{\mu}) = \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j X_i^T X_j$$

$$\text{subject to} \quad 0 \leq \mu_i \leq C, \quad i = 1, \dots, n, \quad \sum_{i=1}^n y_i \mu_i = 0$$

- ▶ The SVM is a linear classifier specified by W^*, b^* :



$$W^* = \sum \mu_i^* y_i X_i, \quad b^* = y_j - X_j^T W^*, \quad (j \text{ s.t. } 0 < \mu_j < C)$$

- ▶ Given a new pattern, X , its class is determined by sign of

$$\begin{aligned} f(X) &= X^T W^* + b^* = \sum_i \mu_i^* y_i X_i^T X + b^* \\ &= \sum_{i: \mu_i > 0} \mu_i^* y_i X_i^T X + (y_j - \sum_{i: \mu_i > 0} \mu_i^* y_i X_i^T X_j) \end{aligned}$$

- ▶ We first formulated SVM for linearly separable case to understand the idea.
- ▶ Then, using slack variables, ξ_i , we have a linear SVM where we make no assumptions about separability.
- ▶ In the dual, the only difference is an upperbound on μ_i .
- ▶ Using this formulation, we can find 'best' hyperplane classifier.
- ▶ The next question is how can we learn non-linear classifiers?
- ▶ Recall that the SVM idea is to transform X_i into some other high-dimensional space and learn a linear classifier there.

Transforming Patterns to become Linearly Separable



Non-linear classifiers

- ▶ In general, we can use a mapping, $\phi : \mathcal{R}^m \rightarrow \mathcal{R}^{m'}$.
- ▶ In $\mathcal{R}^{m'}$, the training set is $\{(Z_i, y_i), i = 1, \dots, n\}$, $Z_i = \phi(X_i)$.
- ▶ We can find optimal hyperplane by solving the dual (replacing $X_i^T X_j$ with $Z_i^T Z_j$).
- ▶ The dual problem now would be the following.

$$\max_{\mu} \quad q(\mu) = \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j \phi(X_i)^T \phi(X_j)$$

$$\text{subject to} \quad 0 \leq \mu_i \leq C, \quad i = 1, \dots, n, \quad \sum_{i=1}^n y_i \mu_i = 0$$

- ▶ The optimization is over \mathcal{R}^n (with quadratic cost function & linear constraints) **irrespective of ϕ and m'** .
- ▶ But computationally expensive?

Kernel function

- ▶ Suppose we have a function, $K : \Re^m \times \Re^m \rightarrow \Re$, that satisfies

$$K(X_i, X_j) = \phi(X_i)^T \phi(X_j)$$

Called Kernel function.

- ▶ Suppose computation of $K(X_i, X_j)$ is about as expensive as that of $X_i^T X_j$.
- ▶ Replacing $\phi(X_i)^T \phi(X_j)$ by $K(X_i, X_j)$, we can solve dual without ever computing any $\phi(X_i)$. Efficient for obtaining optimal hyperplane.
- ▶ What about storing W^* ? Computing $\phi(X)^T W^*$ for new patterns?

Kernel function based classifier

- ▶ Let μ_i^* be soln of Dual. Then $W^* = \sum \mu_i^* y_i \phi(X_i)$.
- ▶ We also have

$$b^* = y_j - \phi(X_j)^T W^* = y_j - \sum_i \mu_i^* y_i \phi(X_i)^T \phi(X_j)$$

- ▶ Given a new pattern X we only need to compute

$$\begin{aligned} f(X) &= \phi(X)^T W^* + b^* \\ &= \sum_i \mu_i^* y_i \phi(X_i)^T \phi(X) + (y_j - \phi(X_j)^T W^*) \\ &= \sum_i \mu_i^* y_i K(X_i, X) + \left(y_j - \sum_i \mu_i^* y_i K(X_i, X_j) \right) \\ &= \sum_{i: \mu_i > 0} \mu_i^* y_i K(X_i, X) + \left(y_j - \sum_{i: \mu_i > 0} \mu_i^* y_i K(X_i, X_j) \right) \end{aligned}$$

- ▶ This is an interesting way of learning nonlinear classifiers.
- ▶ We solve the dual whose dimension is n , number of examples. (We do not need $\phi(X_i)$ for solving dual)
- ▶ All we need to store are:
 - ▶ non-zero Lagrange multipliers: $\mu_i^* > 0$,
 - ▶ Support vectors: X_i, i s.t. $\mu_i^* > 0$.
- ▶ Then, given an X , we compute

$$f(X) = \sum_{i:\mu_i>0} \mu_i^* y_i K(X_i, X) + \left(y_j - \sum_{i:\mu_i>0} \mu_i^* y_i K(X_i, X_j) \right)$$

and classify X based on sign of $f(X)$.

- ▶ Never need to enter ' $\phi(X)$ ' space!

Support Vector Machine

- ▶ Obtain μ_i^* by solving the Dual with $\phi(X_i)^T \phi(X_j)$ replaced by $K(X_i, X_j)$. (Choose a suitable Kernel function. Use 'penalty const', C as needed).
- ▶ Store non-zero μ_i^* and the corresponding support vectors.
- ▶ Classify any new pattern X by sign of

$$f(X) = \sum \mu_i^* y_i K(X_i, X) + \left(y_j - \sum_i \mu_i^* y_i K(X_i, X_j) \right)$$

- ▶ If we have a suitable Kernel function, we never need to compute $\phi(X)$.
- ▶ The range space of ϕ can even be infinite dimensional!

Example kernel function

- ▶ We start with an example kernel function in \Re^2 .
- ▶ Consider $K(X_i, X_j) = (1 + X_i^T X_j)^2$.
- ▶ Let $X_i = (x_{i1}, x_{i2})^T \in \Re^2$ and similarly for X_j .
- ▶ Then

$$K(X_i, X_j) = (1 + x_{i1}x_{j1} + x_{i2}x_{j2})^2$$

- ▶ We now show that there exists a mapping ϕ such that $K(X_i, X_j) = \phi(X_i)^T \phi(X_j)$.

- ▶ Consider $\phi : \Re^2 \rightarrow \Re^6$ given by

$$Z = \phi(X) = [1 \quad x_1^2 \quad x_2^2 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \sqrt{2}x_1x_2]$$

(Here, $X = (x_1 \ x_2) \in \Re^2$).

- ▶ It is easy to see that a linear discriminant function in terms of Z (i.e., in \Re^6) would be a quadratic discriminant function in terms of X (i.e., in \Re^2).
- ▶ Now we show that

$$K(X_i, X_j) = (1 + X_i^T X_j)^2 = Z_i^T Z_j = \phi(X_i)^T \phi(X_j)$$

► Recall

$$Z_i = \phi(X_i) = [1 \quad x_{i1}^2 \quad x_{i2}^2 \quad \sqrt{2}x_{i1} \quad \sqrt{2}x_{i2} \quad \sqrt{2}x_{i1}x_{i2}]$$

We have

$$\begin{aligned} Z_i^T Z_j &= 1 + x_{i1}^2 x_{j1}^2 + x_{i2}^2 x_{j2}^2 + 2x_{i1}x_{j1} + 2x_{i2}x_{j2} + 2x_{i1}x_{i2}x_{j1}x_{j2} \\ &= (1 + x_{i1}x_{j1} + x_{i2}x_{j2})^2 \\ &= (1 + X_i^T X_j)^2 = K(X_i, X_j) \end{aligned}$$

- Easy to see it works for $X \in \Re^n$ in general.
- Thus $K(X_i, X_j) = (1 + X_i^T X_j)^2$ results in a quadratic discriminant function or a quadratic classifier.

- ▶ From this example, it is also easy to see that for a given Kernel function, the mapping ϕ (or the dimension of its range space) is not unique.
- ▶ Consider the same Kernel fn $K(X_i, X_j) = (1 + X_i^T X_j)^2$.
- ▶ Consider the mapping $\phi : \Re^2 \rightarrow \Re^7$ given by

$$Z = \phi(X) = [1 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad x_1^2 \quad x_2^2 \quad x_1x_2 \quad x_1x_2]$$

- ▶ It is easy to see that this mapping also works.

- ▶ We saw that the Kernel $K(X, X') = (1 + X^T X')^2$ results in a quadratic discriminant function (in the original feature space)
- ▶ This is because the effective ϕ function is such that each $x_i x_j$ term is a component of $\phi(X)$.
- ▶ Thus, if $X \in \Re^m$, then any reasonable ϕ function corresponding to this kernel would have range space with dimension $O(m^2)$.
- ▶ Hence, $\phi(X_i)^T \phi(X_j)$ would need $O(m^2)$ multiplications.
- ▶ If we are using a linear SVM, we only need $X_i^T X_j$ which needs m multiplications.
- ▶ When we use the Kernel for the quadratic case, we need only $m + 1$ multiplications.


Example

- ▶ We will first consider a very simple example problem in \Re^2 to get a feel for the method of obtaining SVM.
- ▶ Suppose we have 3 examples:

$$X_1 = (-1, 0), \quad X_2 = (1, 0), \quad X_3 = (0, 0)$$

with $y_1 = y_2 = +1$ and $y_3 = -1$.

Example



- ▶ As is easy to see, a linear classifier is not sufficient here.
- ▶ Suppose we use the Kernel function:
$$K(X, X') = (1 + X^T X')^2.$$

- Recall, the examples are

$$X_1 = (-1, 0), \quad X_2 = (1, 0), \quad X_3 = (0, 0)$$

- The objective function involves $K(X_i, X_j)$. These are given in a matrix below.

$$[(1 + X_i^T X_j)^2] = \begin{bmatrix} 4 & 0 & 1 \\ 0 & 4 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- ▶ The objective function to be maximized is

$$\begin{aligned} q(\boldsymbol{\mu}) &= \sum_{i=1}^3 \mu_i - \frac{1}{2} \sum_{i,j=1}^3 \mu_i \mu_j y_i y_j K(X_i, X_j) \\ &= \sum_{i=1}^3 \mu_i - \frac{1}{2} (4\mu_1^2 + 4\mu_2^2 + \mu_3^2 - 2\mu_1\mu_3 - 2\mu_2\mu_3) \end{aligned}$$

- ▶ The constraints are

$$\mu_1 + \mu_2 - \mu_3 = 0; \quad \text{and} \quad -\mu_i \leq 0, \quad i = 1, 2, 3.$$

- ▶ The lagrangian for this problem is

$$L(\boldsymbol{\mu}, \lambda, \boldsymbol{\alpha}) = q(\boldsymbol{\mu}) + \lambda(\mu_1 + \mu_2 - \mu_3) - \sum_{i=1}^3 \alpha_i \mu_i$$

Here, λ is the Lagrange multiplier for the equality constraint and α_i are the langrange multipliers for the inequality constraints.

- ▶ Using Kuhn-Tucker conditions, we have $\frac{\partial L}{\partial \mu_i} = 0$ and $\mu_1 + \mu_2 - \mu_3 = 0$.
- ▶ This gives us four equations; we have 7 unknowns (Three μ_i , three α_i and λ).
- ▶ By complementary slackness, we have $\alpha_i \mu_i = 0$. Essentially, we need to guess which $\mu_i > 0$.

- ▶ In this simple problem we know all $\mu_i > 0$.
- ▶ This is because all X_i would be support vectors.
- ▶ Hence we take all $\alpha_i = 0$.
- ▶ We have now four unknowns: $\mu_1, \mu_2, \mu_3, \lambda$.
- ▶ Using $\frac{\partial L}{\partial \mu_i} = 0$, $i = 1, 2, 3$ and feasibility, we can solve for μ_i .

► Recall

$$L(\boldsymbol{\mu}, \lambda, \boldsymbol{\alpha}) = q(\boldsymbol{\mu}) + \lambda(\mu_1 + \mu_2 - \mu_3) - \sum_{i=1}^3 \alpha_i \mu_i$$

$$q(\boldsymbol{\mu}) = \sum_{i=1}^3 \mu_i - \frac{1}{2}(4\mu_1^2 + 4\mu_2^2 + \mu_3^2 - 2\mu_1\mu_3 - 2\mu_2\mu_3)$$

- Now $\frac{\partial L}{\partial \mu_i} = 0$, $i = 1, 2, 3$ and feasibility give

$$\begin{aligned}1 - 4\mu_1 + \mu_3 + \lambda &= 0 \\1 - 4\mu_2 + \mu_3 + \lambda &= 0 \\1 - \mu_3 + \mu_1 + \mu_2 - \lambda &= 0 \\\mu_1 + \mu_2 - \mu_3 &= 0\end{aligned}$$

- These give us $\lambda = 1$ and $\mu_3 = 2\mu_1 = 2\mu_2$.
- Thus we get $\mu_1 = \mu_2 = 1$ and $\mu_3 = 2$.
- This completely determines the SVM

- ▶ If we used the penalty constant with $C \geq 2$ we get the same solution.
(If $C < 2$, we can not get this solution).
- ▶ The classification of any X by this SVM is by the sign of $f(X)$:

$$\begin{aligned} f(X) &= \sum_i \mu_i y_i K(X_i, X) + b^* \\ &= K(X_1, X) + K(X_2, X) - 2K(X_3, X) + b^* \end{aligned}$$

- ▶ Let us first calculate b^* .

- Recall the formula

$$b^* = y_j - \sum_i \mu_i y_i K(X_i, X_j), \quad j \text{ s.t. } 0 < \mu_j$$

- Recall

$$[K(X_i, X_j)] = [(1 + X_i^T X_j)^2] = \begin{bmatrix} 4 & 0 & 1 \\ 0 & 4 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- With $j = 1$ we get $b^* = 1 - (4 + 0 - 2) = -1$.
- With $j = 3$ we get $b^* = -1 - (1 + 1 - 2) = -1$.
- If we solved our optimization problem correctly, we should get same b^* !


- ▶ We have $X_1 = (-1, 0)$, $X_2 = (1, 0)$, $X_3 = (0, 0)$ and $K(X, X') = (1 + X^T X')^2$.
- ▶ Hence, taking $X = (x_1, x_2)^T$, we have

$$\begin{aligned} f(X) &= K(X_1, X) + K(X_2, X) - 2K(X_3, X) + b^* \\ &= (1 - x_1)^2 + (1 + x_1)^2 - 2(1) - 1 \\ &= 2x_1^2 - 1 \end{aligned}$$

- ▶ Hence this SVM will assign class +1 to $X = (x_1, x_2)^T$ if

$$2x_1^2 \geq 1 \quad \text{or} \quad |x_1| \geq \frac{1}{\sqrt{2}}$$

Recall Example Patterns



- ▶ Hence this SVM will assign class +1 to $X = (x_1, x_2)^T$ if

$$2x_1^2 \geq 1 \quad \text{or} \quad |x_1| \geq \frac{1}{\sqrt{2}}$$

- ▶ Why not $|x_1| \geq (1/2)$?
- ▶ We are maximizing margin of the hyperplane in ' x^2 '-space.
- ▶ The final SVM is intuitively very reasonable and we solve essentially the same problem whether we are seeking a linear classifier or a nonlinear classifier.

Kernel functions

- ▶ How do we obtain Kernel functions in general?
- ▶ What kind of symmetric functions capture the inner product in an appropriate space?
- ▶ We look at two important characterizations for Kernel functions.

Mercer Kernels

► **Mercer Theorem:**

Given a symmetric function, $K : \mathbb{R}^m \times \mathbb{R}^m \rightarrow \mathbb{R}$,

there exists an inner product space \mathcal{H} and a mapping $\phi : \mathbb{R}^m \rightarrow \mathcal{H}$ so that

$$K(X_1, X_2) = \phi(X_1)^T \phi(X_2)$$

if for all square-integrable functions g ,

$$\int K(X_1, X_2) g(X_1) g(X_2) dX_1 dX_2 \geq 0.$$

Positive definite kernels

- ▶ Let \bar{K} be a $n \times n$ matrix with $\bar{K}_{i,j} = K(X_i, X_j)$.
- ▶ A **positive definite kernel** is the function K such that \bar{K} is positive semi-definite for all n and all data sets $\{X_1, \dots, X_n\}$.
- ▶ That is, given any n , and any feature vectors, X_1, \dots, X_n , we have, for all scalars c_1, \dots, c_n ,

$$\sum_{i,j=1}^n c_i c_j K(X_i, X_j) \geq 0$$

- ▶ If input space is compact, both these notions are same.

- ▶ Now we use Mercer's theorem to show that the function we gave earlier would be a Kernel function.
- ▶ Consider the function

$$K(U, V) = (U^T V)^p = \left(\sum_{i=1}^m u_i v_i \right)^p$$

where $p > 0$ is an integer and

$U = [u_1 \cdots u_m]^T$ and $V = [v_1 \cdots v_m]^T$ are in \Re^m .

- ▶ We want to show that this satisfies the Mercer theorem.

- By expanding the $(U^T V)^p$ we get an expression

$$\left(\sum_{i=1}^m u_i v_i \right)^p = \sum_{r_1, \dots, r_m} \frac{p!}{r_1! r_2! \dots r_m!} \prod_{i=1}^m (u_i v_i)^{r_i}$$

where the summation is over all non-negative integers, r_1, \dots, r_m such that

$$r_1 + r_2 + \dots + r_m = p$$

- We need to show

$$\int_{\mathbb{R}^m} \int_{\mathbb{R}^m} \left(\sum_{i=1}^m u_i v_i \right)^p g(U) g(V) dU dV > 0.$$

- This becomes a sum of integrals by expanding $(\sum u_i v_i)^p$.
- A typical term here is

$$\begin{aligned} & \frac{p!}{r_1! r_2! \dots r_m!} \int \int (u_1 v_1)^{r_1} (u_2 v_2)^{r_2} \dots (u_m v_m)^{r_m} g(U) g(V) dU dV \\ &= \frac{p!}{r_1! r_2! \dots r_m!} \int (u_1)^{r_1} (u_2)^{r_2} \dots (u_m)^{r_m} g(U) dU \\ & \quad \int (v_1)^{r_1} (v_2)^{r_2} \dots (v_m)^{r_m} g(V) dV \\ &= \frac{p!}{r_1! r_2! \dots r_m!} \left(\int u_1^{r_1} u_2^{r_2} \dots u_m^{r_m} g(U) dU \right)^2 \geq 0 \end{aligned}$$

- ▶ Now consider the function

$$K(U, V) = \sum_{j=0}^p a_j (U^T V)^j, \quad a_j \geq 0$$

- ▶ We can show this also satisfies Mercer theorem

$$\begin{aligned} & \int \sum_{j=0}^p a_j (U^T V)^j g(U) g(V) dU dV \\ &= \sum_{j=0}^p a_j \int (U^T V)^j g(U) g(V) dU dV \\ & \geq 0 \end{aligned}$$

- ▶ Hence functions of the form

$$K(X_1, X_2) = \sum_{j=0}^p a_j (X_1^T X_2)^j, \quad a_j \geq 0$$

are kernels (satisfying Mercer's theorem).

- ▶ A special case is

$$K(X_1, X_2) = (1 + X_1^T X_2)^p$$

which is an example we considered earlier.

- ▶ This is called a polynomial kernel.

- ▶ We showed that the function

$$K(X_1, X_2) = (1 + X_1^T X_2)^p$$

satisfies the Mercer theorem.

- ▶ Hence it is a (mercer) kernel.
- ▶ It is easy to show that it satisfies the definition of a positive definite kernel
- ▶ Hence it is also a positive definite kernel

- ▶ Now consider the functions of the type

$$K(U, V) = \sum_{j=0}^{\infty} a_j (U^T V)^j, \quad a_j \geq 0$$

- ▶ Our proof only involved interchanging integration and summation.
- ▶ For finite sum it is always possible.
- ▶ For infinite sum, a sufficient condition is that the above sum is uniformly convergent
- ▶ Then the above would also satisfy Mercer's theorem.

- ▶ Consider the function

$$K(X_1, X_2) = e^{-\frac{(X_1 - X_2)^T (X_1 - X_2)}{2\sigma^2}}$$

- ▶ We can show it satisfies the theorem by noting

$$e^{-(X_1 - X_2)^T (X_1 - X_2)} = e^{-X_1^T X_1} e^{-X_2^T X_2} e^{2X_1^T X_2},$$

and

$$e^{2X_1^T X_2} = \sum_{p=0}^{\infty} \frac{(2X_1^T X_2)^p}{p!}$$

Some Popular Kernel functions

- ▶ Polynomial kernel:

$$K_p(X_1, X_2) = (1 + X_1^T X_2)^p$$

- ▶ Gaussian kernel

$$K_G(X_1, X_2) = e^{-\frac{\|X_1 - X_2\|^2}{\sigma^2}}$$

Generalization abilities

- ▶ SVM idea: learn linear classifier in a transformed space.
- ▶ We said that naively transforming patterns into a high dimensional space does not work.
- ▶ There were two issues in that 'curse of dimensionality'
- ▶ Computational complexity – 'elegantly' solved by the kernel trick.
- ▶ But, does SVM generalize well?
- ▶ We are finding a hyperplane in a very high dimensional space. Do we need very large number of examples?

- ▶ In practice, SVMs perform well.
- ▶ They can learn classifiers that do well on test data without needing (correspondingly) large number of examples.
- ▶ The reason, essentially, is that we learn a hyperplane with large margin.
- ▶ There are different ways to analyze this.
- ▶ We would just state one theoretical result.

Some theoretical results

- ▶ Let P_{err}^n be the error rate on a test set for an SVM trained with n random examples.
- ▶ Then we can show (for SVM with no slack variables)

$$EP_{\text{err}}^n \leq \min \left(\frac{s}{n}, \frac{[R^2 ||W||^2]}{n}, \frac{m}{n} \right)$$

where s is number of support vectors, R is the radius of smallest sphere enclosing all examples, $||W||^{-2}$ is the margin of the maximum margin hyperplane (in the feature space of dimension m).

- ▶ We have

$$EP_{\text{err}}^n \leq \min \left(\frac{s}{n}, \frac{[R^2 ||W||^2]}{n}, \frac{m}{n} \right)$$


- ▶ Optimal hyperplane may generalize well because
 - ▶ 'data compression' is large
 - ▶ Margin is large
 - ▶ dimension of feature space is small

- ▶ The fraction of support vectors is an upperbound on expected generalization error.
- ▶ This is useful as a confidence measure on the learnt SVM.
- ▶ This essentially comes from the concept of 'stability' of a learning algorithm.

- ▶ Let $\{X_1, \dots, X_n\}$ be a data set from which we learnt a classifier.
- ▶ let Z be another random example. If we gave $\{Z, X_2, \dots, X_n\}$ as the training set we expect to learn a very similar classifier.
- ▶ Similarly we can make n data sets by replacing each X_i in turn with Z .
- ▶ The learning algorithm is stable if the errors on the set $\{Z, X_1, \dots, X_n\}$ of all these classifiers are close.
- ▶ One can put a bound on generalization error based on the difference of errors of these classifiers.

Optimal hyperplane and the support Vectors

The optimal hyperplane does not change if we remove any 'non-support-vector' from the training data.




How good is SVM method ?


- ▶ A very competitive method for tackling many PR problems.
- ▶ Learning a nonlinear classifier only involves choosing a Kernel function.
- ▶ User needs to make choice of parameters – kernel function and C . Also parameters of the Opt technique. Bad choices result in 'overfitting'.
- ▶ Support vectors are an important extra information we get.

An example

- Consider the training set for a 2-class problem in \mathbb{R}^2 as below.




- The support vectors are shown below.
(Gaussian Kernel)




Another example

- More complicated 2-class problem (4×4 checker board).



- The support vectors in this example:



Solving the SVM optimization problem

- ▶ So far we have not considered any algorithms for solving for the SVM.
- ▶ We have to solve a constrained optimization problem to obtain the Lagrange multipliers and hence the SVM.
- ▶ Many specialized algorithms have been proposed for this.

- ▶ The optimization problem to be solved is

$$\max_{\boldsymbol{\mu}} \quad q(\boldsymbol{\mu}) = \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j K(X_i, X_j)$$

$$\text{subject to} \quad 0 \leq \mu_i \leq C, \quad i = 1, \dots, n, \quad \sum_{i=1}^n y_i \mu_i = 0$$

- ▶ A quadratic programming (QP) problem with interesting structure.
- ▶ Due to the special structure, many efficient algorithms are proposed.

- ▶ One interesting idea – Chunking
- ▶ We optimize on only a few variables at a time.
- ▶ Dimensionality of the optimization problem is controlled.
- ▶ We keep randomly choosing the subset of variables.
- ▶ Gave rise to the first specialized algorithm for SVM – SVM Light

SMO Algorithm

- ▶ Taking chunking to extreme level – what is the smallest set of variables we can optimize on?
- ▶ We need to consider at least two variables because there is an equality constraint.
- ▶ Sequential Minimal Optimization (SMO) – works on optimizing two variables at a time.
- ▶ We can analytically find the optimum with respect to two variables.
- ▶ The algorithm (heuristically) decides which two we consider in each iteration.
- ▶ A very efficient algorithm
- ▶ There are many such algorithms.

Kernel Trick

- ▶ We use $\phi : \mathbb{R}^n \rightarrow \mathcal{H}$ to map pattern vectors into appropriate high dimensional space.
- ▶ Kernel function allows us to compute innerproducts in \mathcal{H} **implicitly** without using (or even knowing) ϕ .
- ▶ Through kernel functions, we learn nonlinear classifiers using 'linear techniques'.
- ▶ Algorithms that use only innerproducts (e.g., Fisher discriminant, regression etc) can be implicitly executed in a high dimensional, \mathcal{H} , by using a kernel function.
- ▶ We can elegantly construct non-linear versions of linear techniques.

Support Vector Regression

- ▶ Now we consider the regression problem.
- ▶ Given training data

$$\{(X_1, y_1), \dots, (X_n, y_n)\}, \quad X_i \in \mathbb{R}^m, \quad y_i \in \mathbb{R},$$

want to find 'best' function to predict y given X .

- ▶ We search in a parameterized class of functions

$$\begin{aligned} g(X, W) &= w_1 \phi_1(X) + \dots + w_{m'} \phi_{m'}(X) + b \\ &= W^T \Phi(X) + b, \end{aligned}$$

where $\phi_i : \mathbb{R}^m \rightarrow \mathbb{R}$ are some chosen functions.

- ▶ If we choose, $\phi_i(X) = x_i$ (and hence, $m = m'$) then it is the usual linear model.
- ▶ We are essentially learning a linear model in terms of $\Phi(X)$.
- ▶ We want to formulate the problem so that we can use the Kernel idea.
- ▶ Then, by using a kernel function, we never need to compute or even precisely specify the mapping Φ .

Loss function

- ▶ In a general regression problem, we need to find W to minimize

$$\sum_i L(y_i, g(X_i, W))$$

where L is a loss function.

- ▶ We consider a special loss function that allows us to use the kernel trick.

ϵ -insensitive loss

- ▶ We employ ϵ -insensitive loss function:

$$\begin{aligned} L_{\epsilon}(y_i, g(X_i, W)) &= 0 && \text{If } |y_i - g(X_i, W)| < \epsilon \\ &= |y_i - g(X_i, W)| - \epsilon && \text{otherwise} \end{aligned}$$

Here, ϵ is a parameter of the loss function.

- ▶ If prediction is within ϵ of true value, there is no loss.
- ▶ Using absolute value of error rather than square of error allows for better robustness.
- ▶ Also gives us optimization problem with the right structure.
- ▶ Empirical risk minimization under the ϵ -insensitive loss function would minimize

$$\sum_{i=1}^n \max(|y_i - \Phi(X_i)^T W - b| - \epsilon, 0)$$

- ▶ We want W, b to minimize

$$\sum_{i=1}^n \max(|y_i - \Phi(X_i)^T W - b| - \epsilon, 0)$$

- ▶ We can pose the problem as follows.

$$\begin{aligned} \min_{W, b, \xi, \xi'} \quad & \sum_{i=1}^n \xi_i + \sum_{i=1}^n \xi'_i \\ \text{subject to} \quad & y_i - W^T \Phi(X_i) - b \leq \epsilon + \xi_i, \quad i = 1, \dots, n \\ & W^T \Phi(X_i) + b - y_i \leq \epsilon + \xi'_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad \xi'_i \geq 0 \quad i = 1, \dots, n \end{aligned}$$

- ▶ This does not give a dual with the structure we want.
- ▶ So, we reformulate the optimization problem

The Optimization Problem

- Find W, b and ξ_i, ξ'_i to

$$\text{minimize} \quad \frac{1}{2}W^TW + C \left(\sum_{i=1}^n \xi_i + \sum_{i=1}^n \xi'_i \right)$$

$$\begin{aligned} \text{subject to} \quad & y_i - W^T\Phi(X_i) - b \leq \epsilon + \xi_i, \quad i = 1, \dots, n \\ & W^T\Phi(X_i) + b - y_i \leq \epsilon + \xi'_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad \xi'_i \geq 0 \quad i = 1, \dots, n \end{aligned}$$

- We have added the term W^TW in the objective function. This is like model complexity in a regularization context.

- ▶ Like earlier, we can form the Lagrangian and then, using Kuhn-Tucker conditions, can get the optimal values of W and b .
- ▶ Given that this problem is similar to the earlier one, we would get W^* in terms of the optimal lagrange multipliers as earlier.
- ▶ Essentially, the lagrange multipliers corresponding to the inequality constraints on the errors would be the determining factors.
- ▶ We can use the same technique as earlier to formulate the dual to solve for the optimal Lagrange multipliers.

The dual

- ▶ The dual of this problem is

$$\begin{aligned} \max_{\alpha, \alpha'} \quad & \sum_{i=1}^n y_i (\alpha_i - \alpha'_i) - \epsilon \sum_{i=1}^n (\alpha_i + \alpha'_i) \\ & - \frac{1}{2} \sum_{i,j} (\alpha_i - \alpha'_i)(\alpha_j - \alpha'_j) \Phi(X_i)^T \Phi(X_j) \\ \text{subject to} \quad & \sum_{i=1}^n (\alpha_i - \alpha'_i) = 0 \\ & 0 \leq \alpha_i, \alpha'_i \leq C, \quad i = 1, \dots, n \end{aligned}$$

- ▶ Here α_i and α'_i are the Lagrange multipliers corresponding to the first two inequalities in the primal.

The solution

- ▶ We can use the Kuhn-Tucker conditions to derive the final optimal values of W and b as earlier.
- ▶ This gives us

$$\begin{aligned}W^* &= \sum_{i=1}^n (\alpha_i^* - \alpha_i^{*'}) \Phi(X_i) \\b^* &= y_j - \Phi(X_j)^T W^* + \epsilon, \quad j \text{ s.t. } 0 < \alpha_j^* < C/n\end{aligned}$$

- We have

$$W^* = \sum_{i=1}^n (\alpha_i^* - \alpha_i^{*'}) \Phi(X_i)$$

$$b^* = y_j - \Phi(X_j)^T W^* + \epsilon, \quad j \text{ s.t. } 0 < \alpha_j^* < C/n$$

- Note that we have $\alpha_i^* \alpha_i^{*'} = 0$. Also, $\alpha_i^*, \alpha_i^{*'}$ are zero for examples where error is less than ϵ .
- The final W is a linear combination of some of the examples – the support vectors.
- Note that the dual and the final solution are such that we can use the kernel trick.

- ▶ Let $K(X, X') = \Phi(X)^T \Phi(X')$.
- ▶ The optimal model learnt is

$$\begin{aligned} g(X, W^*) &= \sum_{i=1}^n (\alpha_i^* - \alpha_i^{*'}) \phi(X_i)^T \phi(X) + b^* \\ &= \sum_{i=1}^n (\alpha_i^* - \alpha_i^{*'}) K(X_i, X) + b^* \end{aligned}$$

- ▶ As earlier, b^* can also be written in terms of the Kernel function.

Support vector regression

- ▶ Once again, the kernel trick allows us to learn non-linear models using a linear method.
- ▶ The parameters: C , ϵ and parameters of kernel function.
- ▶ The basic idea of SVR can be used in many related problems.

SV regression

- ▶ With the ϵ -insensitive loss function, points whose targets are within ϵ of the prediction do not contribute any 'loss'.
- ▶ Gives rise to some interesting robustness of the method. It can be proved that local movements of target values of points outside the ϵ -tube do not influence the regression.
- ▶ Robustness essentially comes through the support vector representation of the regression.

- ▶ In our formulation of the regression problem we did not explain why we added W^TW term in the objective function.
- ▶ We are essentially minimizing

$$\frac{1}{2}W^TW + C \sum_{i=1}^n \max(|y_i - \Phi(X_i)^TW - b| - \epsilon, 0)$$

- ▶ This is ‘regularized risk minimization’.
- ▶ Then W^TW is the model complexity term which is intended to favour learning of ‘smoother’ models.
- ▶ Next we explain why W^TW is a good term to capture degree of smoothness in case of linear models.

- ▶ Let $f : \mathbb{R}^m \rightarrow \mathbb{R}$ be a continuous function.
- ▶ Continuity means we can make $|f(X) - f(X')|$ as small as we want by taking $\|X - X'\|$ sufficiently small.
- ▶ There are ways to characterize the ‘degree of continuity’ of a function.
- ▶ We consider one such measure now.

ϵ -Margin of a function

- ▶ The ϵ -margin of a function, $f : \Re^n \rightarrow \Re$ is

$$m_\epsilon(f) = \inf\{\|X - X'\| : |f(X) - f(X')| \geq 2\epsilon\}$$

- ▶ The intuitive idea is:

How small can $\|X - X'\|$ be, still keeping $|f(X) - f(X')|$ 'large'

- ▶ The larger $m_\epsilon(f)$, the smoother is the function.

$$m_\epsilon(f) = \inf\{\|X - X'\| : |f(X) - f(X')| \geq 2\epsilon\}$$

- ▶ Obviously, $m_\epsilon(f) = 0$ if f is discontinuous.
- ▶ $m_\epsilon(f)$ can be zero even for continuous functions,
e.g., $f(x) = 1/x$.
- ▶ $m_\epsilon(f) > 0$ for all $\epsilon > 0$ iff f is uniformly continuous.
- ▶ Higher margin would mean the function is 'slowly varying' and hence is a 'smoother' model.

Linear Models and margin

- ▶ Consider regression with linear models. Then,

$$|f(X) - f(X')| = |W^T(X - X')|.$$

- ▶ For all X, X' with $|W^T(X - X')| \geq 2\epsilon$, we want the smallest $\|X - X'\|$
- ▶ It would be smallest if $|W^T(X - X')| = 2\epsilon$ and $(X - X')$ is parallel to W .

That is, $X - X' = \pm \frac{2\epsilon W}{W^T W}$.

- ▶ Thus, $m_\epsilon(f) = \|\pm \frac{2\epsilon W}{W^T W}\| = \frac{2\epsilon}{\|W\|}$.
- ▶ Thus in our optimization problem adding the term $W^T W$ promotes learning of smoother models.
- ▶ As we have seen linear regression models use this as the regularization term.

- ▶ The basic idea of kernel functions, as we saw in SVM, has been extended in many ways.
- ▶ There have been many extensions of the basic SVM method also.
- ▶ Some of them are essentially formulations of approximate solutions to make the algorithm more efficient.
- ▶ Some of them are reformulations to add additional features to the SVM method.
- ▶ We consider a couple of simple examples of such extensions.

- Suppose the optimization problem is changed to

$$\begin{array}{ll}\min_{W, b, \xi} & \frac{1}{2}W^TW + b^2 + C \sum_{i=1}^n \xi_i \\ \text{subject to} & y_i(W^TX_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n\end{array}$$

- We have added the b^2 term to the objective function. The main reason is that it simplifies the dual.

- The dual turns out to be

$$\begin{aligned} \max_{\boldsymbol{\mu}} \quad & \sum_{i=1}^n \mu_i - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j K(X_i, X_j) \\ & - \frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j \end{aligned}$$

subject to $0 \leq \mu_i \leq C, \quad i = 1, \dots, n,$

- The equality constraint is absent.
Only bound constraints on variables.
- Allows for efficient optimization.
(Successive overrelaxation).

- ▶ Next, we consider a reformulation of SVM optimization problem, known as ν -SVM.
- ▶ Recall that the primal problem for SVM with slack variables is

$$\begin{array}{ll} \min_{W, b, \xi} & \frac{1}{2} W^T W + C \sum_{i=1}^n \xi_i \\ \text{subject to} & y_i(W^T \phi(X_i) + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n \end{array}$$

ν -SVM

- ▶ In the SVM formulation with slack variables, we do not know how to choose C .
- ▶ Consider a changed optimization problem

$$\begin{array}{ll}\min_{W, b, \xi, \rho} & \frac{1}{2}W^TW - \nu\rho + \frac{1}{n}\sum \xi_i \\ \text{subject to} & y_i[W^T\phi(X_i) + b] \geq \rho - \xi_i \\ & \xi_i \geq 0.\end{array}$$

where ν is a user-chosen constant.

- ▶ Note that $W, b, \rho, \xi_i = 0$ is a feasible solution.
- ▶ We do not need $\rho \geq 0$ constraint.

- ▶ The Lagrangian for this problem is

$$\begin{aligned} L(W, b, \boldsymbol{\xi}, \rho, \boldsymbol{\eta}, \boldsymbol{\mu}) = & \frac{1}{2}W^TW - \nu\rho + \frac{1}{n}\sum_{i=1}^n \xi_i \\ & - \sum_{i=1}^n \eta_i \xi_i + \sum_{i=1}^n \mu_i (\rho - \xi_i - y_i[W^T\phi(X_i) + b]) \end{aligned}$$

- ▶ The μ_i are the Lagrange multipliers for the separability constraints and η_i are the Lagrange multipliers for the constraints $\xi_i \geq 0$.

The Kuhn-Tucker conditions give us

- ▶ $\nabla_W L = 0 \Rightarrow W = \sum_i \mu_i y_i \phi(X_i)$
- ▶ $\frac{\partial L}{\partial b} = 0 \Rightarrow \sum \mu_i y_i = 0$
- ▶ $\frac{\partial L}{\partial \xi_i} = 0 \Rightarrow \mu_i + \eta_i = \frac{1}{n}, \forall i$
- ▶ $\frac{\partial L}{\partial \rho} = 0 \Rightarrow \sum \mu_i = \nu$
- ▶ $\rho - \xi_i - y_i(W^T \phi(X_i) + b) \leq 0; \quad \xi_i \geq 0; \quad \forall i$
- ▶ $\mu_i \geq 0; \quad \eta_i \geq 0, \quad \forall i$
- ▶ $\mu_i(\rho - \xi_i - y_i(W^T \phi(X_i) + b)) = 0; \quad \eta_i \xi_i = 0, \quad \forall i$

- Suppose $\xi_i > 0$ for some i . Then we have $\eta_i = 0$ and hence $\mu_i = \frac{1}{n}$. Hence

$$\begin{aligned}\nu &= \sum_{i=1}^n \mu_i = \sum_{i: \xi_i > 0} \mu_i + \sum_{i: \xi_i = 0} \mu_i \\ &\geq \sum_{i: \xi_i > 0} \mu_i = \frac{|\{i: \xi_i > 0\}|}{n}\end{aligned}$$

- Hence we have:
 ν is an upper bound on the fraction of ‘margin errors’.

- ▶ We also have, because $0 \leq \mu_i \leq \frac{1}{n}$,

$$\begin{aligned}\nu &= \sum_{i=1}^n \mu_i = \sum_{i: \mu_i > 0} \mu_i + \sum_{i: \mu_i = 0} \mu_i \\ &\leq \sum_{i: \mu_i > 0} \mu_i \leq \frac{|\{i : \mu_i > 0\}|}{n}\end{aligned}$$

- ▶ Hence we have:
 ν is a lower bound on the fraction of support vectors.

- ▶ In the ν -SVM formulation, the ν is the user chosen constant.
- ▶ Unlike the parameter C , the ν has an interesting interpretation.
- ▶ It is simultaneously the upperbound on fraction of errors and lower bound on fraction of support vectors.
- ▶ If for the chosen ν , the problem has a solution with $\rho > 0$, then the bounds would be met.
- ▶ This gives us a good way to choose this 'penalty constant'.

- ▶ The dual for the ν -SVM turns out to be

$$\begin{aligned} \max_{\boldsymbol{\mu}} \quad & q(\boldsymbol{\mu}) = -\frac{1}{2} \sum_{i,j=1}^n \mu_i \mu_j y_i y_j K(X_i, X_j) \\ \text{subject to} \quad & 0 \leq \mu_i \leq \frac{1}{n}, \forall i; \quad \sum_{i=1}^n y_i \mu_i = 0; \quad \sum_{i=1}^n \mu_i = \nu \end{aligned}$$

- ▶ This is a simple optimization problem similar to that of 'C-SVM'.
- ▶ One can show that if we have a solution for ν -SVM then if we choose $C = 1/\rho n$, we get the same solution with 'C-SVM'.

ν SVR

- ▶ This idea can be extended to the regression problem also.
- ▶ In support vector regression, we had two user defined constants: ϵ and C .
- ▶ The ϵ specifies the ‘tolerable error’ and it is difficult to know what value to choose for it.
- ▶ We can reformulate SVR so that we can optimize on ϵ also.
- ▶ This will be very similar to the ν -SVM formulation.

- Recall the optimization problem in SVR:

$$\begin{array}{ll} \min_{W, b, \xi, \xi'} & \frac{1}{2} W^T W + C \left(\sum_{i=1}^n \xi_i + \sum_{i=1}^n \xi'_i \right) \\ \text{subject to} & y_i - W^T \Phi(X_i) - b \leq \epsilon + \xi_i, \quad i = 1, \dots, n \\ & W^T \Phi(X_i) + b - y_i \leq \epsilon + \xi'_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad \xi'_i \geq 0 \quad i = 1, \dots, n \end{array}$$

- We change the optimization problem to the following:

$$\begin{aligned}
 & \min_{W, b, \epsilon, \xi, \xi'} \quad \frac{1}{2} W^T W + C \left(\nu \epsilon + \frac{1}{n} \sum_{i=1}^n (\xi_i + \xi'_i) \right) \\
 & \text{subject to} \quad y_i - W^T \phi(X_i) - b \leq \epsilon + \xi_i, \quad i = 1, \dots, n \\
 & \quad \quad \quad W^T \phi(X_i) + b - y_i \leq \epsilon + \xi'_i, \quad i = 1, \dots, n \\
 & \quad \quad \quad \xi_i \geq 0, \quad \xi'_i \geq 0, \quad \epsilon \geq 0, \quad i = 1, \dots, n
 \end{aligned}$$

where ν is a user-chosen constant.

- We get similar results as in ν -SVM.

Risk minimization view of SVM

- ▶ We posed the support vector regression problem as a (regularized) risk minimization under a special loss function.
- ▶ It was then reformulated into an (equivalent) constrained optimization problem.
- ▶ In contrast, we formulated the SVM directly as a constrained optimization problem.
- ▶ However, it can also be seen to be minimization of (regularized) empirical risk under a special loss function.

- ▶ The optimization problem for SVM is

$$\begin{aligned} \min_{W, b, \xi} \quad & \frac{1}{2} W^T W + C \sum_{i=1}^n \xi_i \\ \text{subject to} \quad & y_i(W^T X_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, n \\ & \xi_i \geq 0, \quad i = 1, \dots, n \end{aligned}$$

- ▶ Given any W, b , the ξ_i have to satisfy

$$\xi_i \geq \max(0, 1 - y_i(W^T X_i + b))$$

- ▶ Since we need to minimize $\sum \xi_i$, we need to take the value above for each ξ_i .

- ▶ Hence we can find SVM by solving the following unconstrained optimization problem:

$$\min_{W,b} \frac{1}{2}W^TW + C \sum_{i=1}^n \max(0, 1 - y_i(W^TX_i + b))$$

- ▶ The model (or classifier) we are learning is $f(X) = W^TX + b$.
- ▶ For this model, we already saw W^TW is a good regularization term.

- ▶ Consider the loss function defined by

$$L_{\text{hinge}}(y, f(X)) = \max(0, 1 - yf(X))$$

- ▶ Then the optimization problem is same as

$$\min_{W,b} \frac{1}{n} \sum_{i=1}^n L(y_i, f(X_i)) + C' \frac{1}{2} W^T W$$

- ▶ Thus, our SVM formulation is empirical risk minimization under hinge-loss along with a regularization term.


- ▶ As we saw earlier, the hinge-loss and square-loss are good convex approximations of the 0–1 loss.
- ▶ For 0–1 loss $L(y, f(X))$ is one if $yf(X)$ is negative and zero otherwise.
- ▶ The squared error loss can be written as

$$L_{\text{square}}(y, h(X)) = (1 - yf(X))^2$$

- ▶ The hinge loss is given by

$$L_{\text{hinge}}(y, h(X)) = \max(0, 1 - yf(X))$$

- We can plot all the functions as follows.



(Here we plot $y f(X)$ on x -axis and $L(y, f(X))$ on y -axis).

- ▶ Hinge loss is also called soft-margin loss.
- ▶ Suppose we want to minimize, over all f ,

$$E[\max(0, 1 - yf(X))], \quad y \in \{+1, -1\}$$

- ▶ Intuitively the best we can do is to make sign of $f(X)$ to be same as sign of the corresponding y .
- ▶ Hence, intuitively, the best f is

$$f(X) > 0, \quad \text{if } P[y = +1|X] > 0.5; \quad \text{else } f(X) < 0$$

- ▶ This is indeed a good classifier.

- ▶ In SVM method, there are two important ingredients.
- ▶ One is the Kernel function.
- ▶ Kernel functions allow us to learn nonlinear models using essentially linear techniques.
- ▶ Second is the 'support vector' expansion – the final model is expressed as a ('sparse') linear combination of some of the data vectors.
- ▶ Kernels are a good way to capture 'similarity' and are useful in general.
- ▶ The support vector expansion is also a general property of Kernel based methods.
- ▶ We look at this general view of Kernels next.